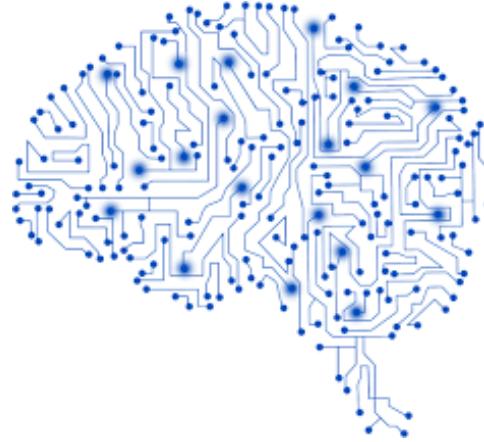
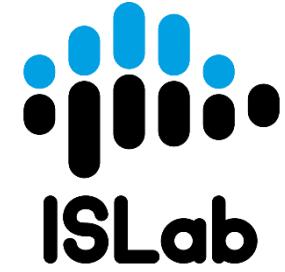


University of Minho
School of Engineering



Deep Learning with TensorFlow™

Connective Systems and Classifiers

Perfil ML:FA @ MiEI/4º ano - 2º Semestre

Bruno Fernandes, Victor Alves

05/03/2020

Contents

2

tf.data API

tf.keras API

TensorBoard

OpenAI

Hands On

- The tf.data API (Mid-level API)
 - tf.data.Dataset
 - Loading and transforming datasets
- The tf.keras API (High-level API)
 - Keras Callbacks e Custom Callbacks
 - Checkpoints, SavedModel and saving Subclassed Models
- TensorBoard
- OpenAI Gym
- Hands On

The tf.data API

3

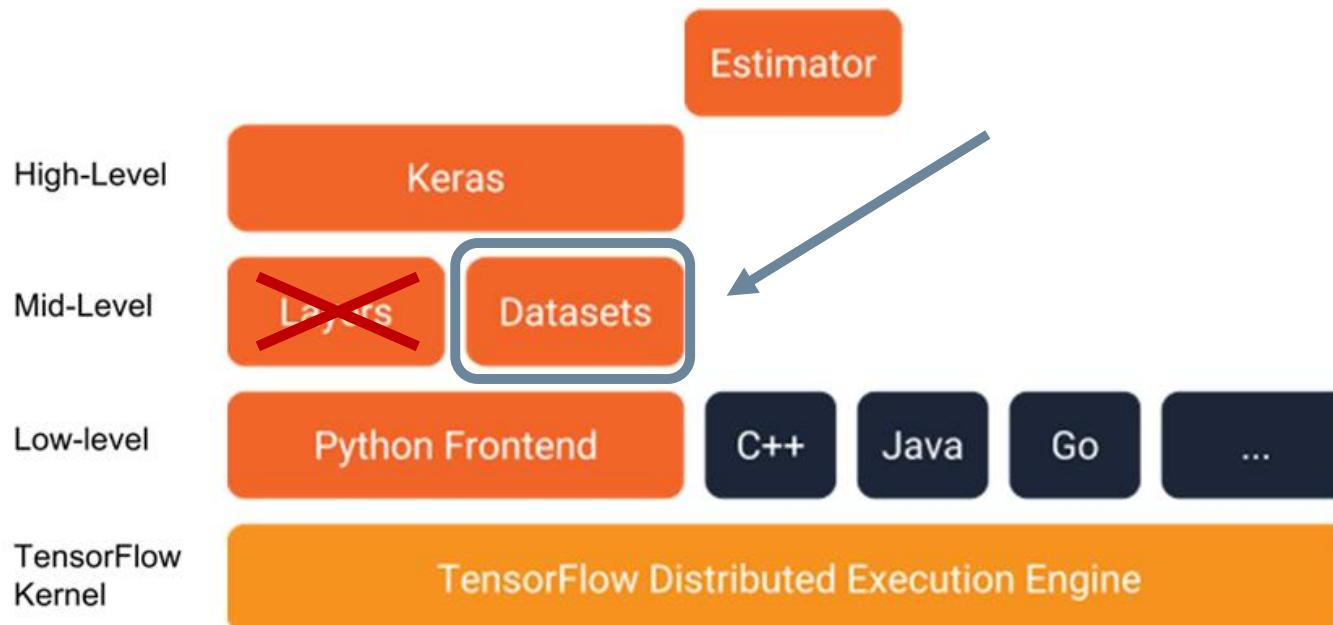
TF.DATA API

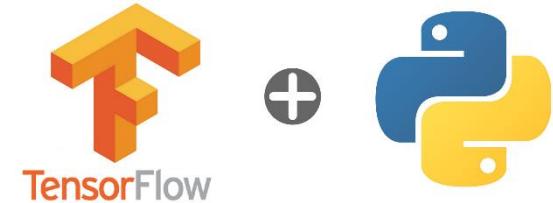
tf.keras API

TensorBoard

OpenAI

Hands On





The tf.data API

4

TF.DATA API

tf.keras API

TensorBoard

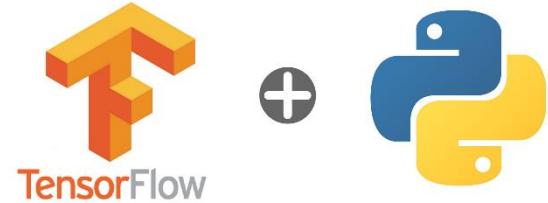
OpenAI

Hands On

TensorFlow introduced the **tf.data API**, a better way to handle data!

- The tf.data API makes it possible to **handle large amounts of data**, read from **different data formats**, and **perform complex transformations**
- The tf.data API introduces a **tf.data.Dataset** abstraction that represents a sequence of elements
- There are two distinct ways to create a dataset:
 - A data **source** builds a Dataset from data stored in **memory** (`tf.data.Dataset.from_tensors` or `tf.data.Dataset.from_tensor_slices`) or in **files** (`tf.data.TFRecordDataset`)
 - A data **transformation** to build a dataset from one or more `tf.data.Dataset` objects

The tf.data.Dataset API



5

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

TensorFlow introduced the **tf.data API**, a better way to handle data!

- The Dataset object is a **Python iterable**. It is possible to consume its elements using a for loop
- A dataset contains elements that each have the same (nested) structure
- The Dataset.element_spec property allows you to inspect the type of each element

```
dataset = tf.data.Dataset.from_tensor_slices([1, 2, 4, 8])  
  
print(dataset)  
print(dataset.element_spec)  
  
for elem in dataset:  
    print(elem)
```

```
<TensorSliceDataset shapes: (), types:  
tf.int32>  
  
TensorSpec(shape=(), dtype=tf.int32,  
name=None)  
  
tf.Tensor(1, shape=(), dtype=int32)  
tf.Tensor(2, shape=(), dtype=int32)  
tf.Tensor(4, shape=(), dtype=int32)  
tf.Tensor(8, shape=(), dtype=int32)
```

The tf.data.Dataset API



6

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

TensorFlow introduced the **tf.data API**, a better way to handle data!

- The Dataset object is a **Python iterable**. It is possible to consume its elements using a for loop
- A dataset contains elements that each have the same (nested) structure
- The Dataset.element_spec property allows you to inspect the type of each element

```
dataset = tf.data.Dataset.from_tensor_slices([1, 2, 4, 8])  
  
print(dataset)  
print(dataset.element_spec)  
  
for elem in dataset:  
    print(elem.numpy())
```

<TensorSliceDataset shapes: (), types: tf.int32>

TensorSpec(shape=(), dtype=tf.int32,
name=None)

1
2
4
8

The tf.data.Dataset API

Load from Numpy



7

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

Features and labels are supposed to be tensors, but since **TensorFlow** and **NumPy** are **seamlessly integrated** they can be NumPy arrays!

Assuming you have a vector of examples and the corresponding labels, pass the **two vectors as a tuple** into `tf.data.Dataset.from_tensor_slices` to create a `tf.data.Dataset`

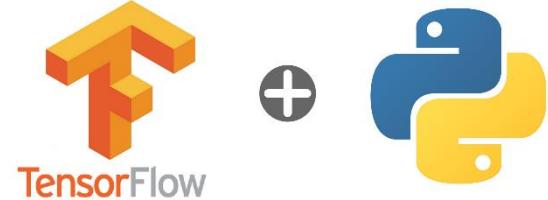
```
path = tf.keras.utils.get_file(FILE_NAME, DATA_URL)

with np.load(path) as data:
    train_examples = data['x_train']
    train_labels = data['y_train']

#creating train dataset from two NumPy arrays
train_dataset = tf.data.Dataset.from_tensor_slices((train_examples, train_labels))
```

The tf.data.Dataset API

Load from pandas.DataFrame



8

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

Or pass the **values of two dataframes** (examples and labels) into `tf.data.Dataset.from_tensor_slices` to create a `tf.data.Dataset`

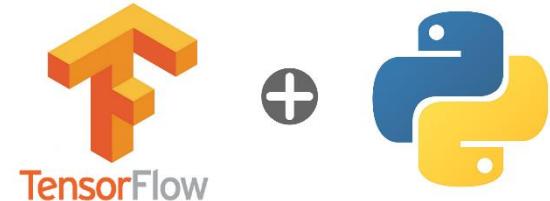
```
path = tf.keras.utils.get_file(FILE_NAME, DATA_URL)

df = pd.read_csv(path)
y = df.pop('class')

#creating train dataset from two dataframes
train_dataset = tf.data.Dataset.from_tensor_slices((df.values, y.values))
```

The tf.data.Dataset API

Load from CSV file



9

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

You can **load csv content using pandas.DataFrame**. But if you need to **scale up to a large set of files** then use the `tf.data.experimental.make_csv_dataset` function.

The only column you need to **identify explicitly** is the `class`.

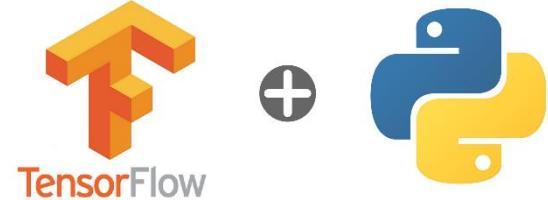
```
#df = pd.read_csv('data.csv')
path = tf.keras.utils.get_file(FILE_NAME, DATA_URL)
#creating raw train dataset from csv file
def get_dataset(file_path, **kwargs):
    dataset = tf.data.experimental.make_csv_dataset(path, batch_size=32, label_name='survived',
        na_value='na', num_epochs=1, ignore_errors=True, **kwargs)
    return dataset

raw_train_dataset = get_dataset(path)

#if the csv file has no column names, pass them as a list of strings to the column_names argument
#instead, use the select_columns argument if you only want a subset of features
CSV_COLUMNS = ['sex', 'age', 'parch', 'fare', 'deck', 'alone', 'survived']
raw_train_dataset = get_dataset(path, column_names=CSV_COLUMNS)
```

The tf.data.Dataset API

Load text file



10

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

Use `tf.data.TextLineDataset` to load examples from **text files**, in which **each example** is a **line of text** from the original file.

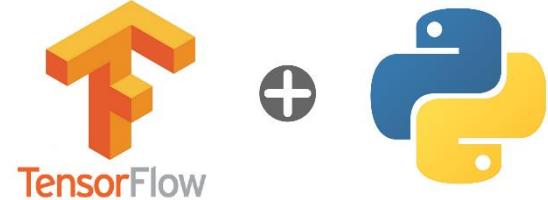
Particularly interesting for NLP.

```
path = tf.keras.utils.get_file(FILE_NAME, DATA_URL)

#creating train dataset from a txt file
lines_dataset = tf.data.TextLineDataset(path)
```

The tf.data.Dataset API

Shuffling & Batching



11

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

Shuffling

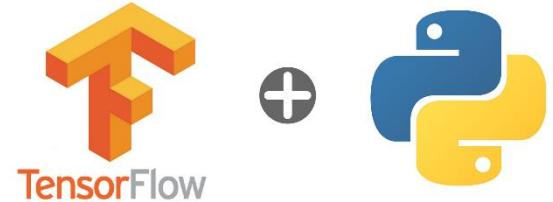
- The Dataset.shuffle() transformation maintains a fixed-size buffer and chooses the next element uniformly at random from that buffer
- Be aware that large buffer_sizes may take a lot of time and memory
- Order matters!

Batching

- The Dataset.batch() transformation **stacks n consecutive elements** of a dataset into a single element
- Dataset.batch results in an unknown batch size because the last batch may not be full!
- Use the **drop_remainder argument** to ignore that last batch, and get full shape

The tf.data.Dataset API

Shuffling & Batching



12

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

```
...
Importing data
...
def import_data():
    #load mnist training and test data
    ➔ (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

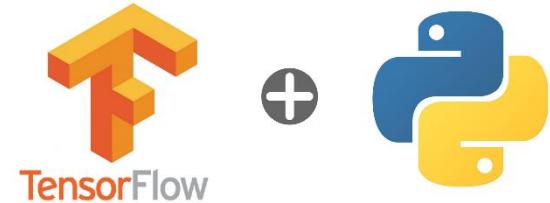
    #some data exploration
    print('***** Log import_data *****')
    print('Train data shape', x_train.shape)
    print('Test data shape', x_test.shape)
    print('Number of training samples', x_train.shape[0])
    print('Number of testing samples', x_test.shape[0])
    for i in range(25):
        plt.subplot(5,5,i+1)      #Add a subplot as 5 x 5
        plt.xticks([])            #get rid of labels
        plt.yticks([])            #get rid of labels
        plt.imshow(x_test[i], cmap="gray")
    plt.show()
    print('***** Log import_data *****')

    #reshape the input to have a list of self.batch_size by 28*28 = 784; and normalize (/255)
    x_train = x_train.reshape(x_train.shape[0], x_train.shape[1]*x_train.shape[2]).astype('float32')/255
    x_test = x_test.reshape(x_test.shape[0], x_test.shape[1]*x_test.shape[2]).astype('float32')/255
    #reserve 5000 samples for validation
    x_validation = x_train[-5000:]
    y_validation = y_train[-5000:]
    #do not use those same 5000 samples for training!
    x_train = x_train[:-5000]
    y_train = y_train[:-5000]

    #create dataset iterator for training
    ➔ train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))
    #shuffle in intervals of 1024 and slice in batches of batch_size
    ➔ train_dataset = train_dataset.shuffle(buffer_size=1024).batch(batch_size)
    #create the validation dataset
    validation_dataset = tf.data.Dataset.from_tensor_slices((x_validation, y_validation))
    validation_dataset = validation_dataset.batch(batch_size)
    return train_dataset, validation_dataset, x_test, y_test
```

The tf.data.Dataset API

Shuffling & Batching



13

TF.DATA API

tf.keras API

TensorBoard

OpenAI

Hands On

```
...
Define a low level fit and predict making use of the tape.gradient
...
def low_level_fit_and_predict():
    #manually, let's iterate over the epochs and fit ourselves
    for epoch in range(epochs):
        print('Epoch %d/%d' %(epoch+1, epochs))

        #to store loss values
        loss_history = []

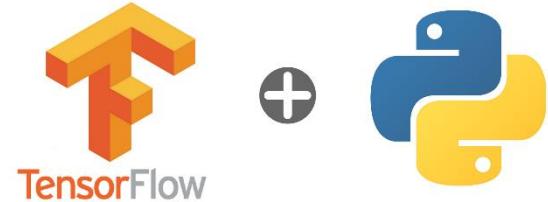
        #iterate over all batchs
        for step, (x_batch, y_batch) in enumerate(train_dataset):
            #use a gradien tape to save computations to calculate gradient later
            with tf.GradientTape() as tape:
                #running the forward pass of all layers
                #operations being recorded into the tape
                probs = mlp(x_batch)
                #computing the loss for this batch
                #how far are we from the correct labels?
                loss_value = loss_object(y_batch, probs)

                #store loss value
                loss_history.append(loss_value.numpy().mean())
                #use the tape to automatically retrieve the gradients of the trainable variables
                #with respect to the loss
                gradients = tape.gradient(loss_value, mlp.trainable_weights)
                #running one step of gradient descent by updating (going backwards now)
                #the value of the trainable variables to minimize the loss
                optimizer.apply_gradients(zip(gradients, mlp.trainable_weights))
                # Update training metric.
                train_metric(y_batch, probs)

                #log every n batches
                if step%200 == 0:
                    print('Step %s. Loss Value = %s; Mean Loss = %s' %(step, str(loss_value.numpy()), np.mean(loss_history)))

            #display metrics at the end of each epoch
            train_accuracy = train_metric.result()
            print('Training accuracy for epoch %d: %s' %(epoch+1, float(train_accuracy)))
            #reset training metrics (at the end of each epoch)
            train_metric.reset_states()
```

tf.data transformations



14

TF.DATA API

tf.keras API

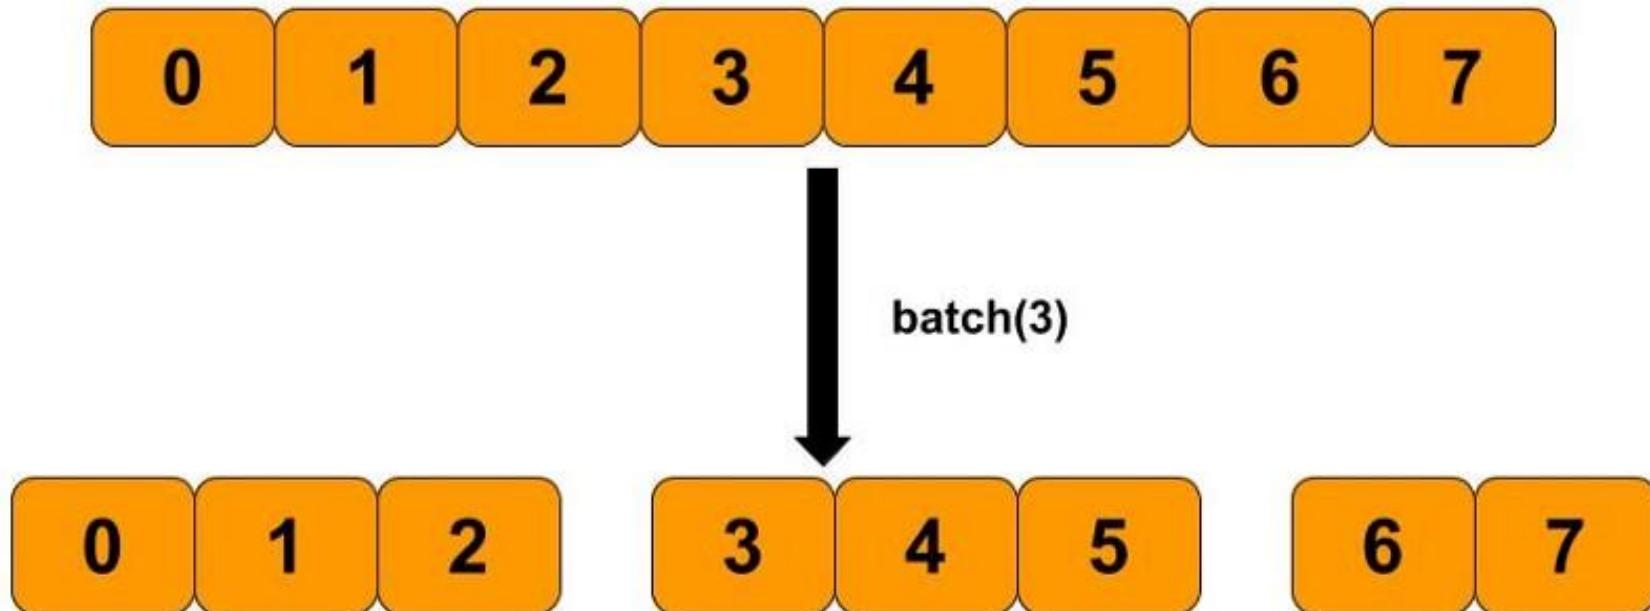
TensorBoard

OpenAI

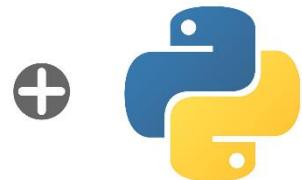
Hands On

Once you have created the Dataset, you can apply several types of transformations (order is important!):

- **Batch** - sequentially dividing your dataset by the specified batch size



tf.data transformations



15

TF.DATA API

tf.keras API

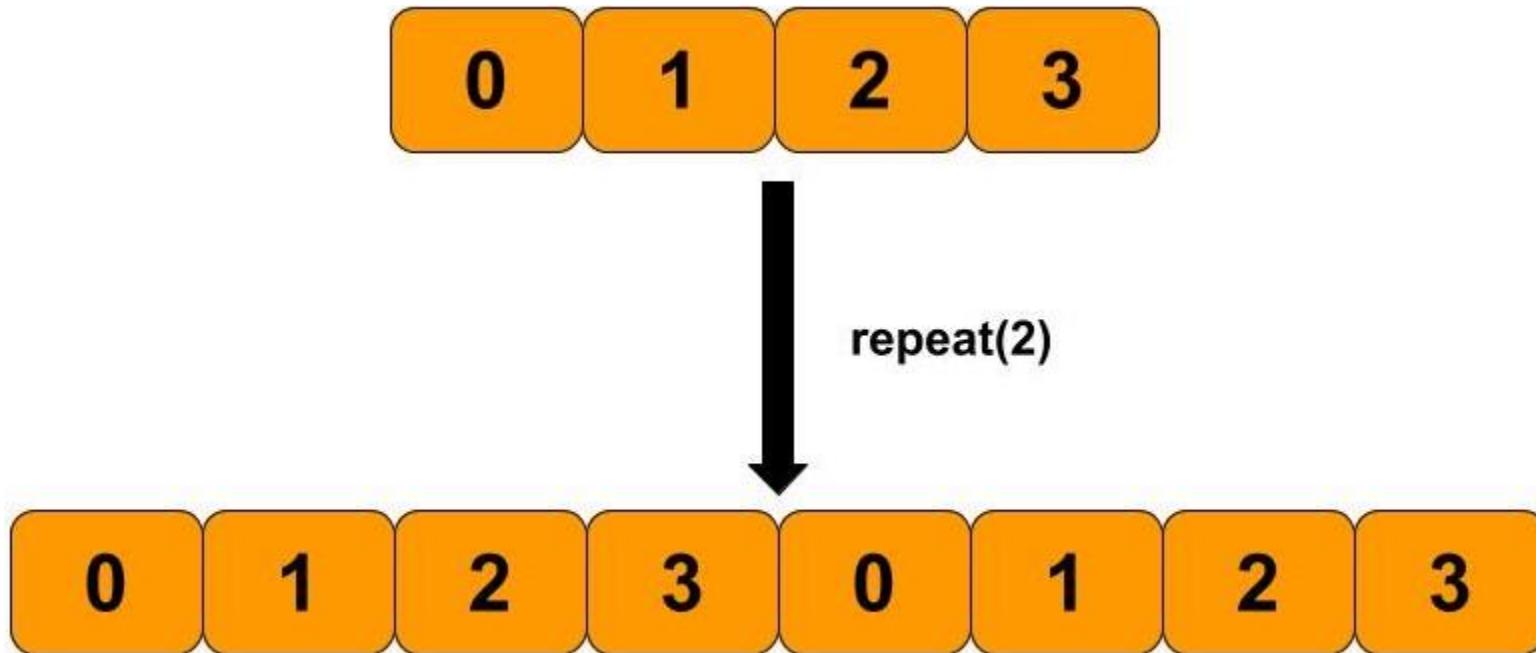
TensorBoard

OpenAI

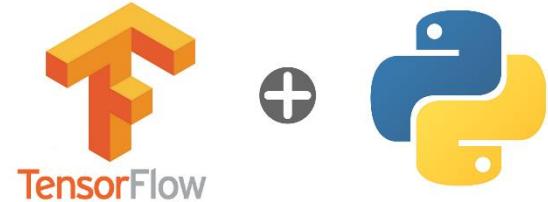
Hands On

Once you have created the Dataset, you can apply several types of transformations (order is important!):

- **Repeat** - use this transformation to create duplicates of the existing data



tf.data transformations



16

TF.DATA API

tf.keras API

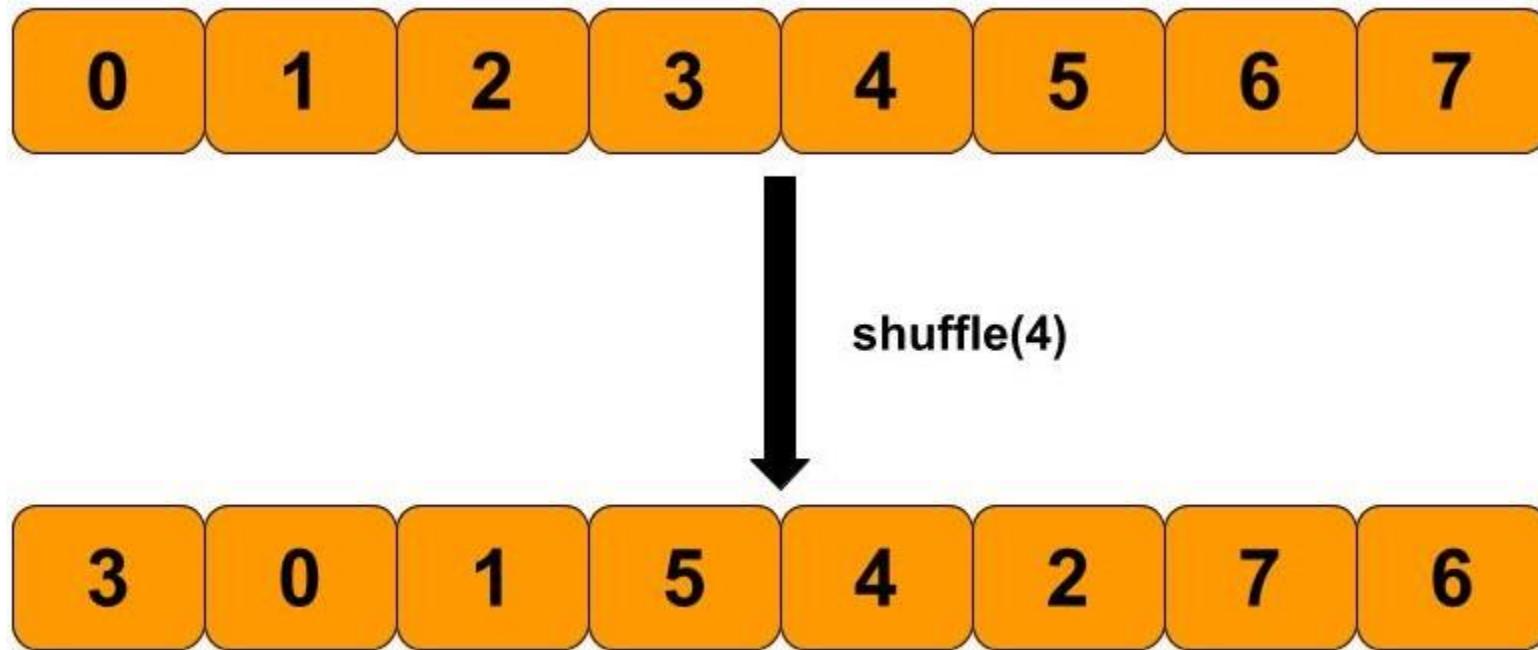
TensorBoard

OpenAI

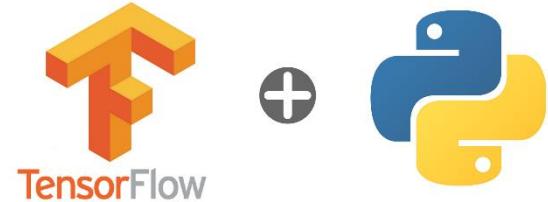
Hands On

Once you have created the Dataset, you can apply several types of transformations (order is important!):

- **Shuffle** - randomly shuffle the data in the Dataset



tf.data transformations



17

TF.DATA API

tf.keras API

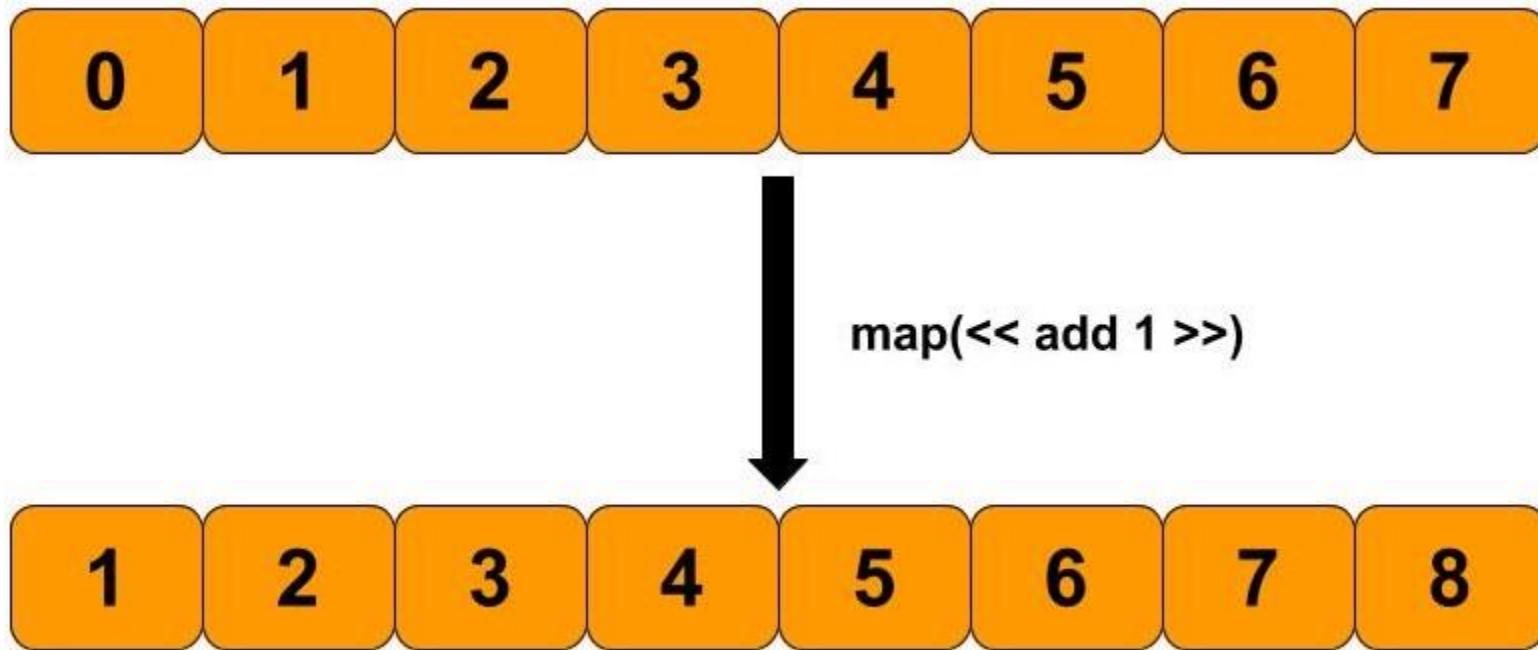
TensorBoard

OpenAI

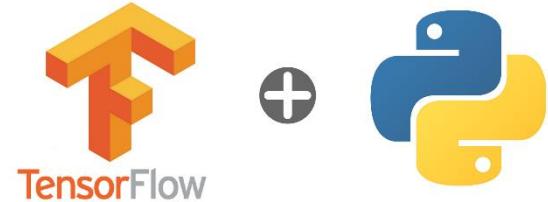
Hands On

Once you have created the Dataset, you can apply several types of transformations (order is important!):

- **Map** - apply some operation to all the individual data elements in the dataset



tf.data transformations



18

TF.DATA API

tf.keras API

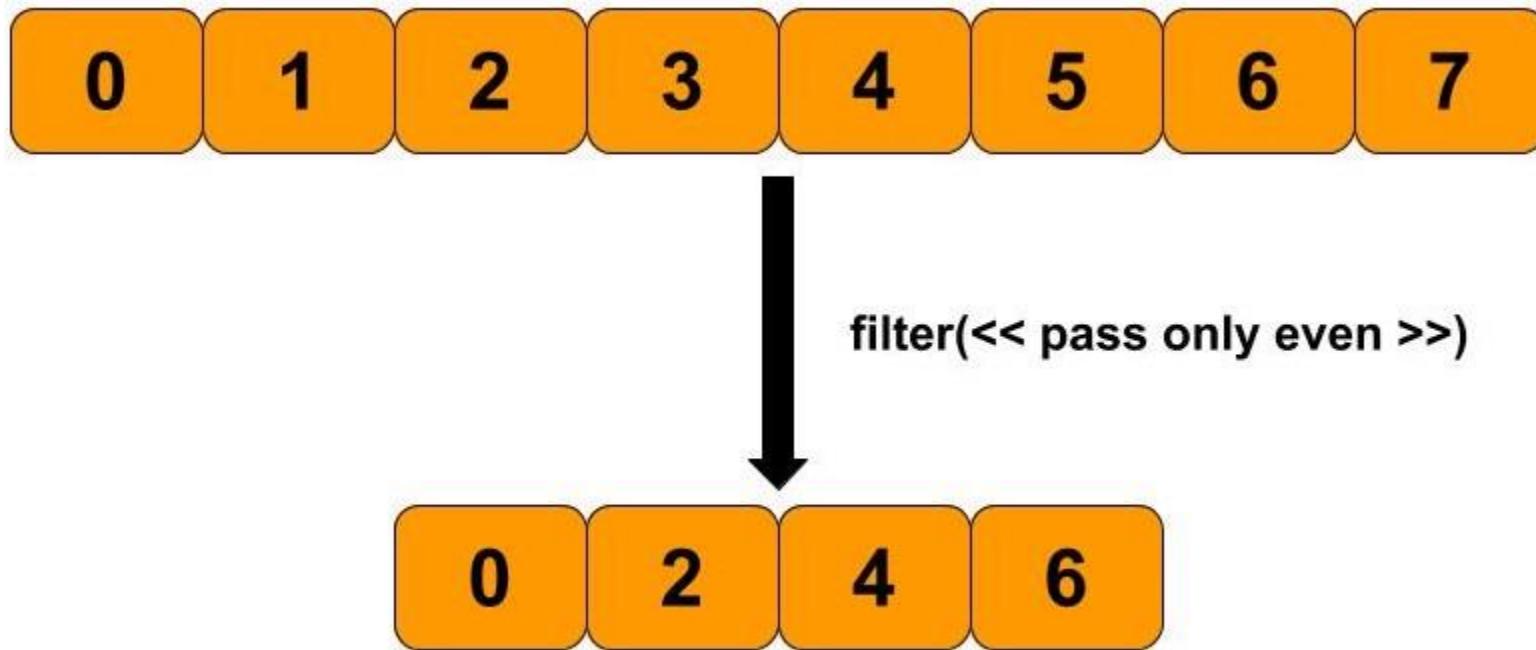
TensorBoard

OpenAI

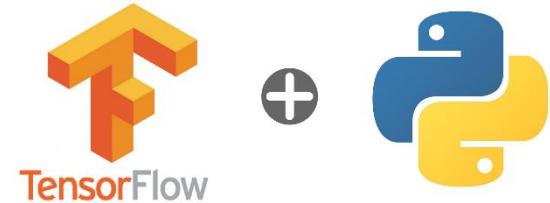
Hands On

Once you have created the Dataset, you can apply several types of transformations (order is important!):

- **Filter** - filter out some elements from the Dataset



High-level TensorFlow APIs



19

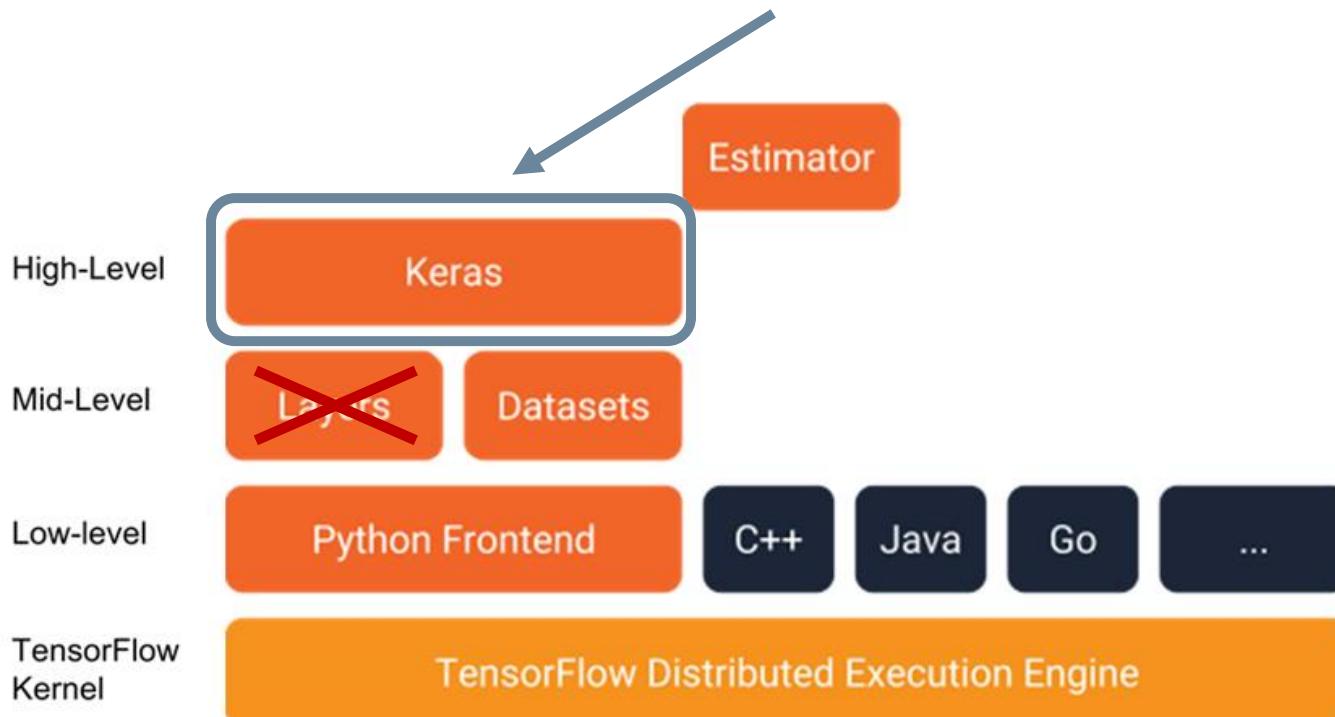
tf.data API

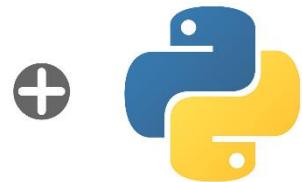
TF.KERAS API

TensorBoard

OpenAI

Hands On





20

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

Keras: The Python Deep Learning library

Keras is a **high-level neural networks API**, written in Python and capable of running on top of TensorFlow. Indeed, Keras is now (since v2.0) the **official high-level API of TensorFlow!**

Use Keras if you need a deep learning library that:

- Allows for **easy** and **fast prototyping** (through user friendliness, modularity, and extensibility)
- Supports both **convolutional networks** and **recurrent networks**, as well as combinations of the two
- Runs seamlessly on **CPU** and **GPU**

Officially, Keras is compatible with **Python 2.7 to 3.6!**



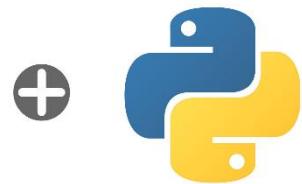
Keras

21

tf.data API

TF.KERAS API

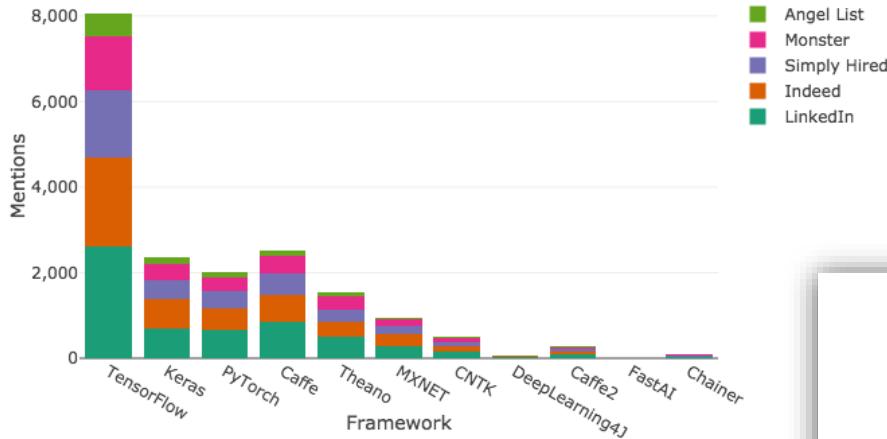
TensorBoard



OpenAI

Hands On

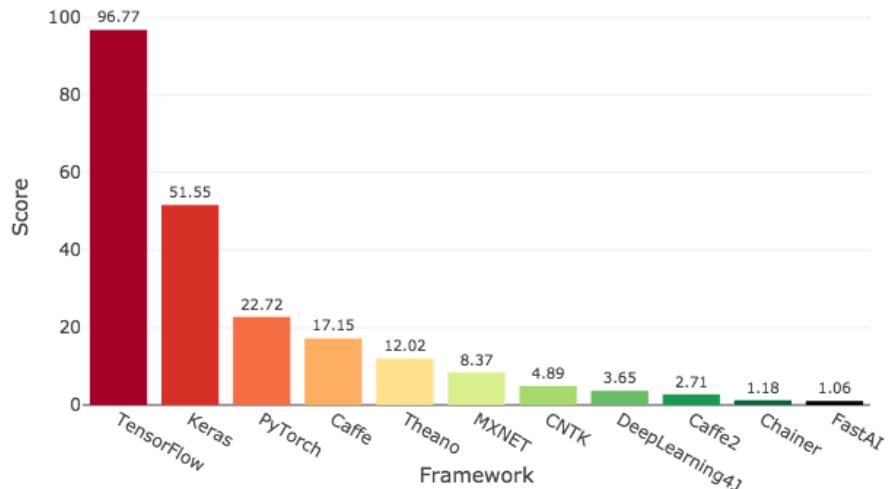
Online Job Listings



What's special about Keras?

- Focused on user experience
- Large adoption in the industry and research community
- Multi-backend, multi-platform
- Easy productization of models

Deep Learning Framework Power Scores 2018





Keras

22

tf.data API

TF.KERAS API

TensorBoard



OpenAI

Hands On

NETFLIX

UBER

Google



Get Started



23

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

tf.keras is TensorFlow's implementation of the [Keras API specification](#). It is a high-level API to build and train models that includes support for TensorFlow-specific functionality!

tf.keras makes TensorFlow easier to use without sacrificing flexibility and performance!

To get started, **import tf.keras** as part of your TensorFlow program as such:

```
import tensorflow as tf  
  
print(tf.__version__)  
  
print(tf.keras.__version__)
```

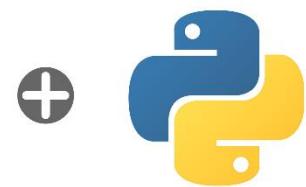


2.0.0

2.2.4-tf

Three API styles

K Keras +



24

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

The Sequential Model

- As simple as it gets
- Used for single-input, single-output and sequential layer stacks
- Good for >70% of use cases

The functional API

- Like playing with Lego bricks
- Multi-input, multi-output, arbitrary static graph topologies
- Good for 95% of use cases

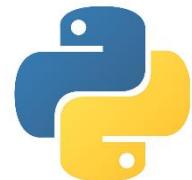
Model subclassing

- Maximum flexibility
- Larger potential for errors

Sequential model



Keras



25

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

The core data structure of Keras is a **model**, a way to **organize layers**. The simplest and most common type of model is the **Sequential model**, a linear stack of layers.

To build a simple, fully-connected network (i.e. multi-layer perceptron):

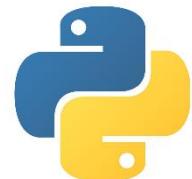
```
import tensorflow as tf
from tensorflow.keras import layers

#sequential model
model = tf.keras.Sequential()
#flatten layer to transform input from 2d (28x28) to 1d (784)
model.add(layers.Flatten(input_shape=(28, 28)))
# add a densely-connected layer with 64 neurons
model.add(layers.Dense(64, activation='relu'))
#add a softmax layer with 10 output neurons
model.add(layers.Dense(10, activation='softmax'))
```

Sequential model



Keras



26

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

The core data structure of Keras is a **model**, a way to **organize layers**. The simplest and most common type of model is the **Sequential model**, a linear stack of layers.

To build a simple, fully-connected network (i.e. multi-layer perceptron):

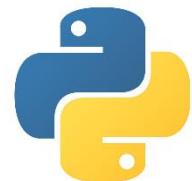
```
import tensorflow as tf
from tensorflow.keras import layers

#sequential model
model = tf.keras.Sequential([
    layers.Flatten(input_shape=(28, 28)),
    layers.Dense(64 , activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

Sequential model



Keras



27

tf.data API

TF.KERAS API

TensorBoard

OpenAI

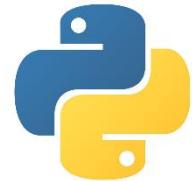
Hands On

There are many `tf.keras.layers` available with some common constructor parameters:

- `activation` - set the activation function for the layer (defaults to no activation being applied)
- `kernel_initializer` and `bias_initializer` - the initialization schemes that create the layer's weights (defaults to Glorot/Xavier uniform initializer)
- `kernel_regularizer` and `bias_regularizer` - apply penalties on layer parameters or layer activity during optimization (penalties are incorporated in the loss function)

```
#create a sigmoid activated layer
layers.Dense(64, activation='sigmoid')
#linear layer with L1 regularization (Lasso Regression) of factor 0.01 applied to weights
layers.Dense(64, kernel_regularizer=tf.keras.regularizers.l1(0.01))
#linear layer with L2 regularization (Ridge Regression) of factor 0.01 applied to the bias
layers.Dense(64, bias_regularizer=tf.keras.regularizers.l2(0.01))
#linear layer with a bias vector initialized to 2.0s
layers.Dense(64, bias_initializer=tf.keras.initializers.constant(2.0))
```

Compile it



28

tf.data API

TF.KERAS API

TensorBoard

OpenAI

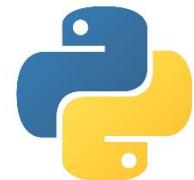
Hands On

After building the model, **configure the model for training** by calling the **compile** method

```
#sequential model
model = tf.keras.Sequential()
#flatten layer to transform input from 2d (28x28) to 1d (784)
model.add(layers.Flatten(input_shape=(28, 28)))
# add a densely-connected layer with 64 neurons
model.add(layers.Dense(64, activation='relu'))
#add a softmax layer with 10 output neurons
model.add(layers.Dense(10, activation='softmax'))

#configure the model
model.compile(optimizer=tf.train.AdamOptimizer(0.001),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

Compile it



29

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

`tf.keras.Model.compile` takes three important arguments:

- `optimizer` - specifies the training procedure. Pass optimizer instances from `tf.keras.optimizers` module
- `loss` - function to minimize during optimization. Loss functions are specified by name or by passing a callable object from the `tf.keras.losses` module
- `metrics` - used to monitor training. May be strings or callable objects from the `tf.keras.metrics` module

```
#configure the model
model.compile(optimizer=tf.train.AdamOptimizer(0.001),
              loss='mse',
              metrics=['mae'])

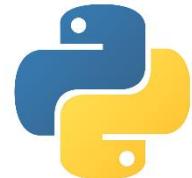
#or

model.compile(optimizer=tf.train.GradientDescentOptimizer(0.01),
              loss=tf.keras.losses.categorical_crossentropy,
              metrics=[tf.keras.metrics.categorical_accuracy])
```

Fit the model



Keras



30

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

`tf.keras.Model.fit` also takes three important arguments:

- `epochs` - training is structured into epochs. An epoch is one iteration over the entire input data (normally done in smaller batches)
- `batch_size` - specifies the size of each batch
- `validation_data` - data on which to evaluate the loss and any model metrics at the end of each epoch. The model will not be trained on this data

```
#training
model.fit(data, labels, epochs=10, batch_size=32)

#or (the tf.data.dataset was already shuffled and batched)
model.fit(train_dataset, validation_data=validation_dataset, epochs=epochs)
```

NumPy inputs



31

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

For **small datasets**, use **in-memory NumPy arrays** to train and evaluate the model

```
import numpy as np

#numpy data
data = np.random.random((1000, 32))
labels = np.random.random((1000, 10))

#validation data
val_data = np.random.random((100, 32))
val_labels = np.random.random((100, 10))

# ... building the model

#train the model
model.fit(data, labels, epochs=10, batch_size=32,
           validation_data=(val_data, val_labels))
```

To scale to **large datasets** use the tf.data API

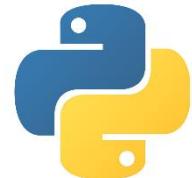
```
#instantiate a dataset instance
dataset = tf.data.Dataset.from_tensor_slices((data, labels))
dataset = dataset.batch(32)
dataset = dataset.repeat(2)

...
#specify steps_per_epoch when calling fit on a dataset
model.fit(dataset, epochs=10, steps_per_epoch=30)
```

Note:

- **steps_per_epoch** refers to the **total number of steps before declaring one epoch finished** and starting the next epoch. When training with TensorFlow data tensors, the default is equal to the number of samples in the dataset divided by the batch size, or 1 if that cannot be determined

Evaluate & Predict K Keras



33

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

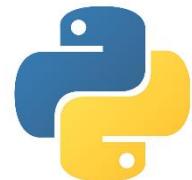
To **evaluate loss** and **metrics** for the provided data:

```
import numpy as np
...
#numpy data
test_data = np.random.random((1000, 32))
test_labels = np.random.random((1000, 10))
...
test_loss, test_acc = model.evaluate(test_data, test_labels, batch_size=32)
#or
test_loss, test_acc = model.evaluate(dataset, steps=30)
```

And to **predict** the output of the last layer in inference for the provided data:

```
predictions = model.predict(pred_data, batch_size=32)
```

Functional API



34

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

A more flexible way of creating models than the Sequential API.

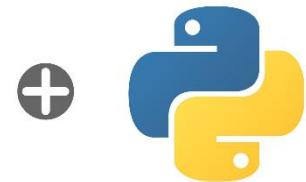
Use the Functional API to **build complex model topologies** such as:

- Multi-input models
- Multi-output models
- Models with shared layers (the same layer called several times)

Building a model with the **functional API** works like this:

1. A layer instance is callable and returns a tensor
2. Input tensors and output tensors are used to define a `tf.keras.Model` instance
3. This model is trained just like the Sequential model.

Functional API



35

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

```
#returns a placeholder input
inputs = tf.keras.Input(shape=(784,))

#a layer instance is callable on a tensor and returns a tensor.
x = tf.keras.layers.Dense(64, activation='relu')(inputs)
#or
#dense = layers.Dense(64, activation='relu')
#x = dense(inputs)

x = tf.keras.layers.Dense(64, activation='relu', name='second_layer')(x)
probs = tf.keras.layers.Dense(10, activation='softmax')(x)

#instantiate the model given the inputs and outputs
model = tf.keras.Model(inputs=inputs, outputs=probs, name='mnist')

print(model.summary())
#you must install pydot and graphviz to plot your model
tf.keras.utils.plot_model(model, 'my_mnist_model.png', show_shapes=True)

#call model.compile(), model.fit() and model.evaluate() as you do with the Sequential API
```

Functional API



36

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

Model: "mnist"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 784)]	0
dense (Dense)	(None, 64)	50240
second_layer (Dense)	(None, 64)	4160
dense_1 (Dense)	(None, 10)	650
=====		

Total params: 55,050

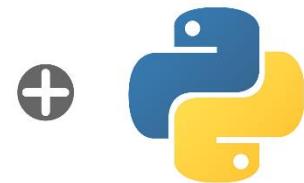
Trainable params: 55,050

Non-trainable params: 0

None

Functional API

K Keras +



37

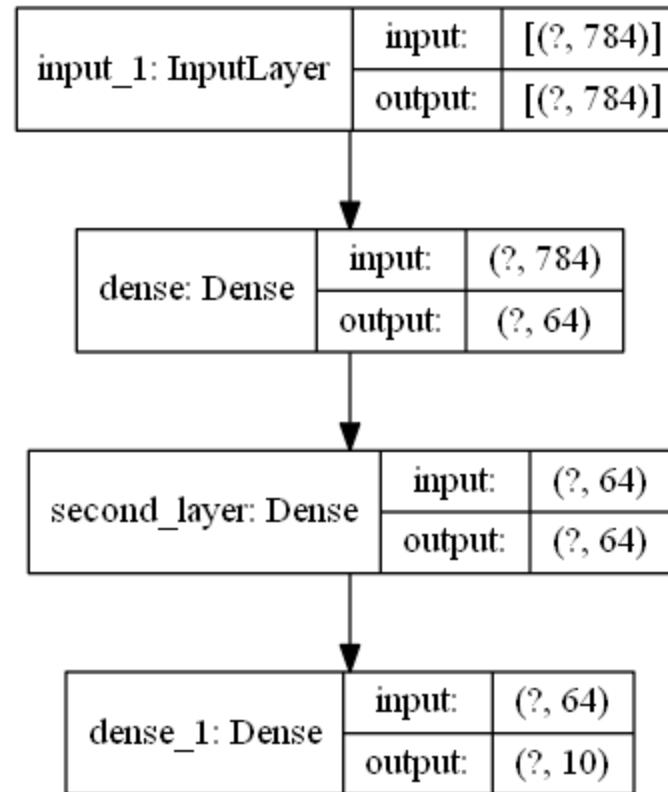
tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

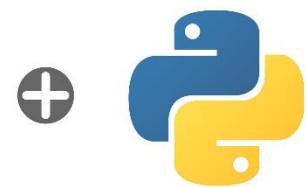


Allows us to build a **fully-customizable model** by subclassing `tf.keras.Model` and **defining our own forward pass!**

We can also create our custom layer by subclassing `tf.keras.layers.Layer`!

Where have we seen this??

Model Subclassing



39

tf.data API

TF.KERAS API

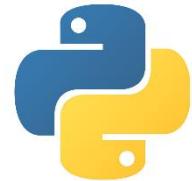
TensorBoard

OpenAI

Hands On

```
# =====
# PerceptronLayer Class
#
# =====
class PerceptronLayer(Layer):
    ...
    The constructor in an object oriented perspective.
    Called when an object is created, allowing the class to initialize the attributes of a class.
    neurons corresponds to the number of neurons in this perceptron layer
    ...
    def __init__(self, neurons=16, **kwargs):
        super(PerceptronLayer, self).__init__(**kwargs)
        self.neurons = neurons
    ...
    We use the build function to deferr weight creation until the shape of the inputs is known
    ...
    def build(self, input_shape):
        pass
    ...
    Implements the function call operator (when an instance is used as a function).
    It will automatically run build the first time it is called, i.e., layer's weights are created dynamically
    ...
    def call(self, inputs):
        pass
    ...
    Enable serialization on our perceptron layer
    ...
    def get_config(self):
        pass
```

Model Subclassing



40

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

```
# =====
#       Multilayer Perceptron Class
#
# =====
class MultilayerPerceptron(Model):
    ...
    The Layers of our MLP (with a fixed number of neurons)
    ...
    def __init__(self, output_neurons=10, name='multilayerPerceptron', **kwargs):
        super(MultilayerPerceptron, self).__init__(name=name, **kwargs)
        self.perceptron_layer_1 = PerceptronLayer(16)
        self.perceptron_layer_2 = PerceptronLayer(32)
        self.perceptron_layer_3 = PerceptronLayer(output_neurons)

    ...
    Layers are recursively composable, i.e.,
    if you assign a Layer instance as attribute of another Layer, the outer layer will start tracking the weights of the inner layer.
    Remember that the build of each layer is called automatically (thus creating the weights).
    ...
    def feed_model(self, input_data):
        pass

    """
    Compute softmax values for the logits
    """
    def softmax(self, logits):
        pass

    def print_trainable_weights(self):
        pass

    def call(self, input_data):
        pass
```

Model Subclassing



41

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

...

Layers are recursively composable, i.e., if you assign a Layer instance as attribute of another Layer, the outer layer will start tracking the weights of the inner layer. Remember that the build of each layer is called automatically (thus creating the weights).

```
...
def feed_model(self, input_data):
    x = self.perceptron_layer_1(input_data)
    #activation function applied to the output of the perceptron layer
    x = tf.nn.relu(x)
    #the output, now normalized, is fed as input to the second perceptron layer
    x = self.perceptron_layer_2(x)
    #again, activation function applied to the output of the second perceptron layer
    x = tf.nn.relu(x)
    #which, again, is fed as input to the third layer, which returns its output
    logits = self.perceptron_layer_3(x)
    #the output of the last layer going over a softmax activation
    #so, we will not be outputting logits but "probabilities"
    return self.softmax(logits) #equivalent of tf.nn.softmax(logits)
```

Using the model

High Level



42

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

```
...  
Define a high level fit and predict making use tf.Keras APIs  
...  
  
def high_level_fit_and_predict():  
    #shortcut to compile and fit a model!  
    #able to do this because our model subclasses tf.keras.Model  
    mlp.compile(optimizer, loss=loss_object, metrics=[train_metric])  
    #since the train_dataset already takes care of batching, we don't pass a batch_size argument  
    #passing validation data for monitoring validation loss and metrics at the end of each epoch  
    history = mlp.fit(train_dataset, validation_data=validation_dataset, epochs=epochs)  
    #print('\nHistory values per epoch:', history.history)  
  
    #evaluating the model on the test data  
    print('\nEvaluating model on test data...')  
    scores = mlp.evaluate(x_test, y_test, batch_size=batch_size, verbose=0)  
    print('Evaluation %s: %s' %(mlp.metrics_names, str(scores)))  
  
    #finally, generating predictions (the output of the last layer)  
    print('\nGenerating predictions for ten samples...')  
    predictions = mlp.predict(x_test[:10])  
    #now, for each prediction in predictions, get the value with higher "probability"  
    #look at the shape, it is as (3, 10). For each prediction, we have the prob of beeing 0, beeing 1, etc...  
    #we now choose the index of the list with higher "probability"  
    #if pos=3 is the one with higher probability it means it predicts a 3  
    print('Predictions shape:', predictions.shape)  
    for i, prediction in enumerate(predictions):  
        #tf.argmax returns the INDEX with the largest value across axes of a tensor  
        predicted_value = tf.argmax(prediction)  
        label = y_test[i]  
        print('Predicted a %d. Real value is %d.' %(predicted_value, label))
```

Using the model

Low Level



43

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

```
'''  
Define a Low Level fit and predict making use of the tape.gradient  
...  
  
def low_level_fit_and_predict():  
    #manually, let's iterate over the epochs and fit ourselves  
    for epoch in range(epochs):  
        print('Epoch %d/%d' %(epoch+1, epochs))  
  
        #to store loss values  
        loss_history = []  
  
        #iterate over all batchs  
        for step, (x_batch, y_batch) in enumerate(train_dataset):  
            #use a gradient tape to save computations to calculate gradient later  
            with tf.GradientTape() as tape:  
                #running the forward pass of all layers  
                #operations being recorded into the tape  
                probs = mlp(x_batch)  
                #computing the loss for this batch  
                #how far are we from the correct labels?  
                loss_value = loss_object(y_batch, probs)  
  
                #store loss value  
                loss_history.append(loss_value.numpy().mean())  
                #use the tape to automatically retrieve the gradients of the trainable variables  
                #with respect to the loss  
                gradients = tape.gradient(loss_value, mlp.trainable_weights)  
                #running one step of gradient descent by updating (going backwards now)  
                #the value of the trainable variables to minimize the loss  
                optimizer.apply_gradients(zip(gradients, mlp.trainable_weights))  
                # Update training metric.  
                train_metric(y_batch, probs)  
  
                #log every n batches  
                if step%200 == 0:  
                    print('Step %s. Loss Value = %s; Mean Loss = %s' %(step, str(loss_value.numpy()), np.mean(loss_history)))  
  
            #display metrics at the end of each epoch  
            train_accuracy = train_metric.result()  
            print('Training accuracy for epoch %d: %s' %(epoch+1, float(train_accuracy)))  
            #reset training metrics (at the end of each epoch)  
            train_metric.reset_states()
```

Callbacks



44

tf.data API

TF.KERAS API

TensorBoard

OpenAI

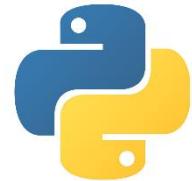
Hands On

A **callback** is an **object passed to a model to customize** and **extend its behavior** during training:

- `tf.keras.callbacks.LearningRateScheduler` - dynamically change the learning rate
- `tf.keras.callbacks.EarlyStopping` - interrupt training when validation performance has stopped improving
- `tf.keras.callbacks.ModelCheckpoint` - model is automatically saved during training
- `tf.keras.callbacks.TensorBoard` - write TensorBoard logs

```
#defining callbacks
callbacks = [
    #interrupt training if loss stops improving for over 2 epochs
    tf.keras.callbacks.EarlyStopping(patience=2, monitor='cost'),
    #write TensorBoard logs to ./tf_logs directory
    tf.keras.callbacks.TensorBoard(log_dir='./logs')
]
model.fit(data, labels, batch_size=32, epochs=5, callbacks=callbacks,...)
```

Custom Callbacks



45

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

We can write our own **custom callback** too!

```
#extend tf.keras.callbacks.Callback
class OurCustomCallback(tf.keras.callbacks.Callback):

    #logs dict contains the loss value and all the metrics at the end of a batch/epoch
    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch %d begins.', %batch)

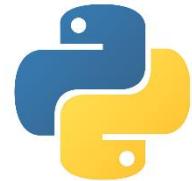
    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch %d ends.', %batch)

    def on_test_batch_begin(self, batch, logs=None):
        print('Evaluating: batch %d begins.', %batch)

    def on_test_batch_end(self, batch, logs=None):
        print('Evaluating: batch %d ends.', %batch)

model.fit(data, labels, batch_size=32, epochs=5, callbacks=[OurCustomCallback()])
```

Custom Callbacks



46

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

We can write our own **custom callback** too!

Override:

- `on_(train|test|predict)_begin(self, logs=None)`
 - Called at the beginning/end of fit/evaluate/predict.
- `on_(train|test|predict)_batch_begin(self, batch, logs=None)`
 - Called right before/after processing a batch during training/testing/predicting.
- `on_epoch_(begin|end)(self, epoch, logs=None)`
 - For training only

Checkpoints



47

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

There may be times where we need to **let our models run for days...** And then:

- We decide to adjust something...!!
- The computer crashes!
- Power goes off!
- A memory error emerges!
- ... the list goes on...

Training is interrupted and we will need to run our model all over again!!!



Checkpoints



48

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

It is crucial to have the ability to **stop training at any point**, for any reason, and **resume training later** as if nothing happened!

- **Checkpoints** capture the exact value of all `tf.Variable` objects used by a model
- Checkpoints **do not contain any description of the computation** defined by the model
- Typically **useful** when **source code** of the original model is **available**

Tip:

- A good practice is to periodically save the model's parameters so that we can later restore or retrain our model from a previous step, if required.

Checkpoints



49

tf.data API

TF.KERAS API

TensorBoard

OpenAI

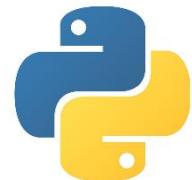
Hands On

We can use the `tf.keras.callbacks.ModelCheckpoint` callback to automatically save checkpoints during training

```
#defining callbacks
callbacks = [
    tf.keras.callbacks.ModelCheckpoint(
        #saving in Keras HDF5 (or h5), a binary data format
        filepath='ckpt/my_model_{epoch}_{val_loss:.3f}.hdf5', #path where to save the model
        save_best_only=True, #overwrite the current checkpoint if and only if
        monitor='val_loss', #the val_loss score has improved
        save_weights_only=False, #if True, only the weights are saved
        verbose=1) #verbosity mode
]

#if the filepath does not contain formatting options like {epoch} or {val_loss}
#then filepath will be overwritten by each new better model
model.fit(data, labels, batch_size=32, epochs=5, callbacks=callbacks,...)
```

Checkpoints



50

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

In the filesystem...

```
C:\ Linha de comandos
Directory of C:\data\PythonWorkspace\mlfa_csc_tf2\ckpt

30/01/2020  15:11    <DIR>        .
30/01/2020  15:11    <DIR>        ..
30/01/2020  15:11            71 552 my_model_10_0.152.hdf5
30/01/2020  15:11            71 552 my_model_15_0.151.hdf5
30/01/2020  15:10            71 552 my_model_1_0.294.hdf5
30/01/2020  15:10            71 552 my_model_2_0.242.hdf5
30/01/2020  15:11            71 552 my_model_3_0.215.hdf5
30/01/2020  15:11            71 552 my_model_4_0.192.hdf5
30/01/2020  15:11            71 552 my_model_5_0.183.hdf5
30/01/2020  15:11            71 552 my_model_6_0.172.hdf5
30/01/2020  15:11            71 552 my_model_7_0.165.hdf5
30/01/2020  15:11            71 552 my_model_8_0.155.hdf5
30/01/2020  15:11            71 552 my_model_9_0.155.hdf5
               11 File(s)   787 072 bytes
                2 Dir(s)  33 199 054 848 bytes free

C:\data\PythonWorkspace\mlfa_csc_tf2\ckpt>
```

Hands On Checkpoints



51

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

Implement this callback in our low-level MLP (@ function `high_level_fit_and_predict`)!

Note: you will need to set the name for each weight and bias!
(issue <https://github.com/tensorflow/tensorflow/issues/26811>)

```
#defining callbacks
callbacks = [
    tf.keras.callbacks.ModelCheckpoint(
        #saving in Keras HDF5 (or h5), a binary data format
        filepath='ckpt/my_model_{epoch}_{val_loss:.3f}.hdf5', #path where to save the model
        save_best_only=True, #overwrite the current checkpoint if it's better
        monitor='val_loss', #the val_loss score has improved
        save_weights_only=False, #if True, only the weights are saved
        verbose=1) #verbosity mode
]

#if the filepath does not contain formatting options like {epoch} or {val_loss}
#then filepath will be overwritten by each new better model
model.fit(data, labels, batch_size=32, epochs=5, callbacks=callbacks,...)
```

HANDS ON

`tf.keras.Model.save()` function includes:

- The model's **architecture**
- The model's **weight values** (learned during training)
- The model's **training config** (passed to the `compile` function)
- The **optimizer** and its state, if any (enables you to restart training)

```
#save the model
model.save('my_model.h5')

#recreate the exact same model from file
new_model = tf.keras.models.load_model('my_model.h5')
new_predictions = new_model.predict(x_test)
```

However if we apply it to our sub-classed, low-level MLP...

NotImplementedError: Saving the model to HDF5 format requires the model to be a Functional model or a Sequential model.

It does not work for subclassed models, because such models are defined via the body of a Python method, which isn't safely serializable.

Consider saving to the Tensorflow SavedModel format (by setting `save_format="tf"`) or using ``save_weights``.

The `tf.keras.Model.save()` function works only for **Functional or Sequential models!**

SavedModel



54

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

We can also export a whole model to the TensorFlow **SavedModel** format:

- SavedModel is a **standalone serialization format** for TensorFlow objects
- It contains a **complete TensorFlow program**, including weights and computation
- Models in this format are **independent** of the source code that created the model, which makes it useful for sharing and deploying

SavedModel



55

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

The SavedModel files include:

- A **TensorFlow checkpoint** containing the model weights
- A **SavedModel proto** containing the underlying TensorFlow graph

The **SavedModel** also works for **Functional** or **Sequential** models!

```
#save a savedmodel
model.save('my_model', save_format='tf')

#recreate the exact same savedmodel from file
new_model = tf.keras.models.load_model('my_model')
new_predictions = new_model.predict(x_test)
```

Saving Subclassed Models



56

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

- Sequential and Functional models are data structures that represent a Directed Acyclic Graph of layers that **can be safely serialized** and **deserialized**
- On the other hand, a **subclassed model** is a piece of code, where the architecture of the model is defined via the body of the call method
- To **load a subclassed model** we need to have **access to the code** that created it
- Also, a **subclassed model that has never been used/trained cannot be saved!** Remember that a subclassed model needs to be called on some data in order to build the weights (otherwise the shape and dtype of the input data is unknown)!

Saving Subclassed Models



57

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

Hence, the recommended way to **save a subclassed model** is to use **save_weights function**, creating a SavedModel checkpoint that contains the value of all variables:

- Layer's Weights
- Optimizer's state
- Other variables associated with stateful model metrics

Some notes on **loading a subclassed model**:

- To restore the optimizer state and the state of any stateful metric you should compile the model with the exact same arguments as before
- Then call it on some data
- Then call `load_weights`

Saving Subclassed Models



58

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

```
#save a savedmodel
model.save_weights('my_weights', save_format='tf')
---

#to restore the optimizer state and the state of any stateful metric
#you should compile the model with the exact same arguments as before
new_model = MultilayerPerceptron(output_neurons=output_neurons)
new_model.compile(optimizer, loss=loss_object, metrics=[train_metric])

#and call it on some data before calling load_weights
#initialize the variables used by the optimizer as well as any stateful metric variables
new_model.train_on_batch(x_train[:1], y_train[:1])

#load the weights of the old model
new_model.load_weights('my_weights')

new_predictions = new_model.predict(x_test)
```

Saving to json



59

tf.data API

TF.KERAS API

TensorBoard

OpenAI

Hands On

Saving the **model architecture** to **json** as well as the **weights** of a trained model. However, it will **not include the training configuration** and the **optimizer**. You will need to call `compile()` again before using the model.

```
#serialize a model to JSON format
json_config = model.to_json()
with open('model_config.json', 'w') as json_file:
    json_file.write(json_config)
#save weights
model.save_weights('my_weights.h5') #Keras HDF5 format
#model.save_weights('my_tf_checkpoint', save_format='tf') #tf checkpoint format

#load json model from the 2 files
with open('model_config.json') as json_file:
    json_config = json_file.read()
new_model = tf.keras.models.model_from_json(json_config)
new_model.load_weights('my_weights.h5')
```

TensorBoard



Keras



60

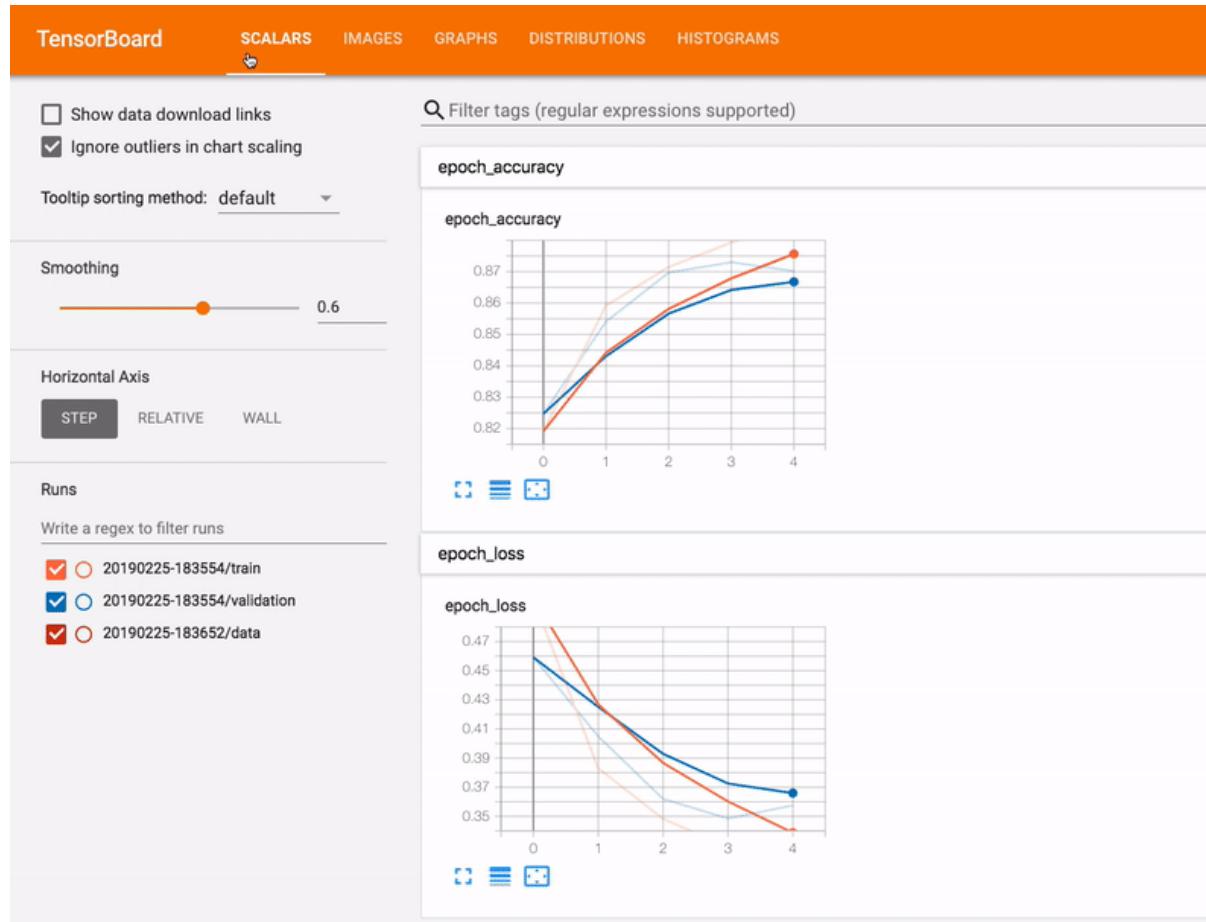
tf.data API

tf.keras API

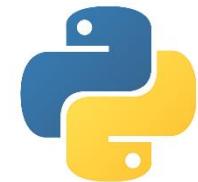
TENSORBOARD

OpenAI

Hands On



What is it?



61

tf.data API

tf.keras API

TENSORBOARD

OpenAI

Hands On

TensorFlow's visualization toolkit:

- The computations we will be using TensorFlow for can be complex and confusing
- TensorBoard is a **set of visualization tools** to make it easier to **understand, debug and optimize TensorFlow programs**
- TensorBoard can be used to visualize our TensorFlow graph, plot quantitative metrics about the execution of the graph or show additional data
- The graph visualizer is a component of TensorBoard that renders the structure of your graph visually in a browser

Why?



62

tf.data API

tf.keras API

TENSORBOARD

OpenAI

Hands On

TensorBoard provides the **visualization** and **tooling** needed for ML experimentation:

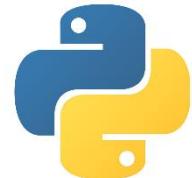
- **Tracking** and **visualizing metrics** such as loss and accuracy
- Visualizing the **model graph** (ops and layers)
- Viewing **histograms** of weights, biases, or other tensors as they change over time
- Projecting embeddings to a lower dimensional space
- Displaying images, text, and audio data
- Profiling TensorFlow programs
- ...

However, it tends to be more complex than it should...

Using TensorBoard



Keras



63

tf.data API

tf.keras API

TENSORBOARD

OpenAI

Hands On

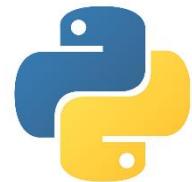
```
#load the TensorBoard notebook extension (run it in the IPython shell)
%load_ext tensorboard
```

```
#a simple ANN
import tensorflow as tf
import datetime
import os
#mnist dataset
mnist = tf.keras.datasets.mnist
#data normalization
(x_train, y_train),(x_test, y_test) = mnist.load_data()
x_train, x_test = x_train/255.0, x_test/255.0
#model creation
def create_model():
    return tf.keras.models.Sequential([
        tf.keras.layers.Flatten(input_shape=(28, 28)),
        tf.keras.layers.Dense(512, activation='relu'),
        tf.keras.layers.Dropout(0.2),
        tf.keras.layers.Dense(10, activation='softmax')
    ])
```

Using TensorBoard



Keras



64

tf.data API

tf.keras API

TENSORBOARD

OpenAI

Hands On

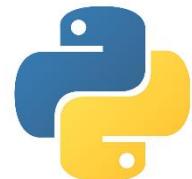
```
model = create_model()
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

#placing the logs in a timestamped subdirectory
log_dir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))
#this callback is used to ensure that logs are created and stored
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

#fit the model using the training data, for 5 epochs, passing the previous callback
model.fit(x=x_train,
           y=y_train,
           epochs=5,
           validation_data=(x_test, y_test),
           callbacks=[tensorboard_callback])
```

```
#start TensorBoard through the notebook (works nicely with jupyter and colab)
%tensorboard --logdir logs
#in anaconda prompt/terminal run as follows (if you use spyder or pycharm)
tensorboard --logdir logs --port 7070
```

Running it



65

tf.data API

tf.keras API

TENSORBOARD

OpenAI

Hands On

After executing our program, go to the terminal (or Anaconda Prompt for windows users) and run:

A screenshot of an Anaconda Prompt window. The title bar says "Anaconda Prompt". The command line shows the user activating a conda environment and then running TensorBoard. The output indicates that TensorBoard is serving at port 7070.

```
(base) C:\Users\bruno>conda activate mlfa_csc_tf2  
(mlfa_csc_tf2) C:\Users\bruno>tensorboard --logdir logs --port 7070  
Serving TensorBoard on localhost; to expose to the network, use a proxy or pass --bind_all  
TensorBoard 2.0.0 at http://localhost:7070/ (Press CTRL+C to quit)
```

Running it



66

tf.data API

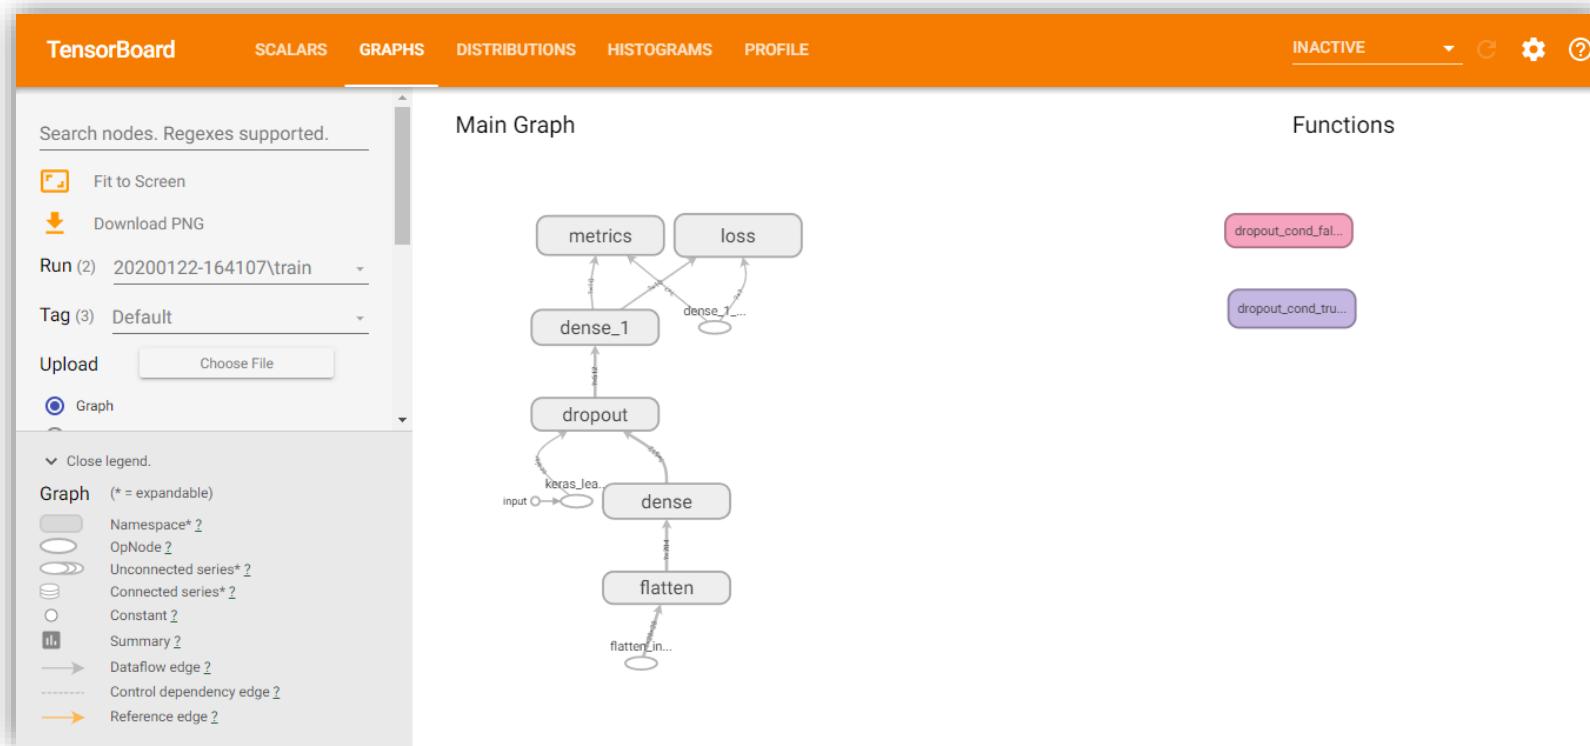
tf.keras API

TENSORBOARD

OpenAI

Hands On

Then, just open your browser and go to <http://localhost:7070/>
(change to the selected port)



OpenAI Gym

K Keras +



67

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

Environments Documentation



Gym

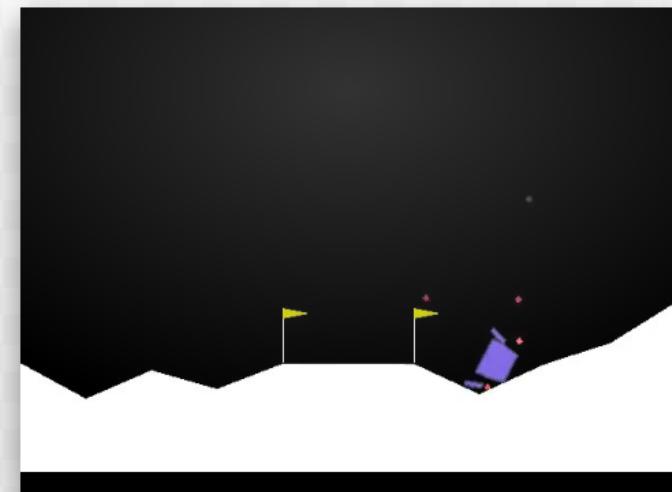
Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball.

[View documentation >](#)

[View on GitHub >](#)



RandomAgent on SpaceInvaders-v0



RandomAgent on LunarLander-v2

Environments Documentation



Algorithms

Atari

Box2D

Classic control

MuJoCo

Robotics NEW

Toy text EASY

Classic control

Control theory problems from the classic RL literature.



[Acrobot-v1](#)

Swing up a two-link robot.



CartPole-v1
Balance a pole on a cart.



MountainCar-v0
Drive up a big hill.

Environments Documentation 

Algorithms

Atari
Reach high scores in Atari 2600 games.

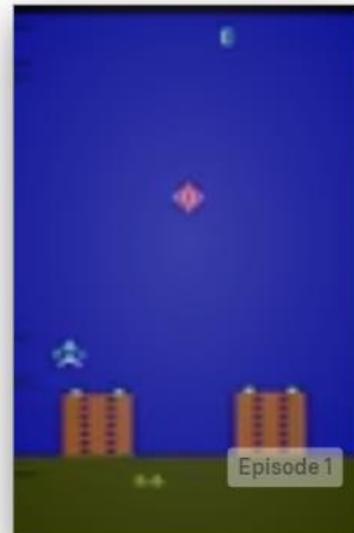
Box2D

Classic control

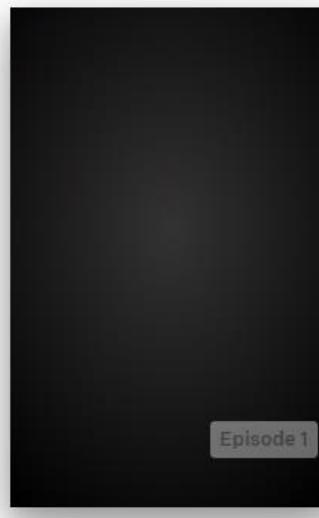
MuJoCo

Robotics NEW

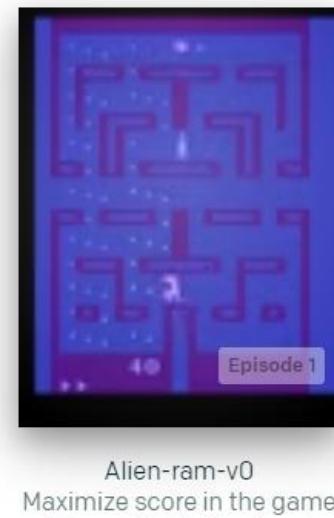
Toy text EASY



AirRaid-ram-v0



AirRaid-v0



Alien-ram-v0
Maximize score in the game Alien, with RAM as input

Environments Documentation



Getting Started with Gym

Gym is a toolkit for developing and comparing reinforcement learning algorithms. It makes no assumptions about the structure of your agent, and is compatible with any numerical computation library, such as TensorFlow or Theano.

The `gym` library is a collection of test problems — **environments** — that you can use to work out your reinforcement learning algorithms. These environments have a shared interface, allowing you to write general algorithms.

Installation

To get started, you'll need to have Python 3.5+ installed. Simply install `gym` using `pip`:

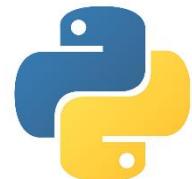
```
pip install gym
```

And you're good to go!

Building from Source

If you prefer, you can clone along the [GitHub repository directly](#). This is particularly useful when working

OpenAI Gym



71

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

openai / gym

Watch 933 Unstar 15,608 Fork 4,076

Code Issues 357 Pull requests 64 Projects 0 Wiki Insights

README.rst

Status: Maintenance (expect bug fixes and minor updates)

OpenAI Gym

OpenAI Gym is a toolkit for developing and comparing reinforcement learning algorithms. This is the `gym` open-source library, which gives you access to a standardized set of environments.

[build passing](#)

See What's New section below

`gym` makes no assumptions about the structure of your agent, and is compatible with any numerical computation library, such as TensorFlow or Theano. You can use it from Python code, and soon from other languages.

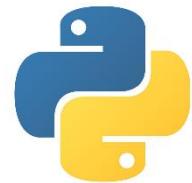
If you're not sure where to start, we recommend beginning with the [docs](#) on our site. See also the [FAQ](#).

A whitepaper for OpenAI Gym is available at <http://arxiv.org/abs/1606.01540>, and here's a BibTeX entry that you can use to cite it in a publication:

```
@misc{1606.01540,  
Author = {Greg Brockman and Vicki Cheung and Ludwig Pettersson and Jonas Schneider and John Schulman and Jie Tang and Wojciech Zaremba and Peter Abbeel and Ilya Sutskever},  
Title = {OpenAI Gym},  
Year = {2016}}
```

OpenAI Gym

The CartPole-v1 Example



72

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

The screenshot shows the Spyder Python IDE interface. On the left, the code editor displays `TP1_example.py` with Python code for a Player Agent. The code initializes a game environment, defines a `PlayerAgent` class with methods like `__init__` and `play_random_games`, and handles rendering and logging. In the center, a small window shows the `CartPole-v1` environment where a pole is balanced on a cart. On the right, the IPython console shows the training log, which lists epochs and batches with their respective losses.

```
15 # -----
16 #
17 #           Player Agent
18 #
19 #
20 class PlayerAgent:
21     ...
22     Constructor - defining game variables
23     ...
24
25     def __init__(self, game_steps, nr_of_games_for_training):
26         self.game_steps = game_steps
27         self.nr_of_games_for_training = nr_of_games_for_training
28         self.score_requirement = score_requirement
29         self.random_games = 50
30         self.log = log
31         self.env = gym.make(game)
32         self.scores = []
33
34     ...
35     Let us test the environment and play some games!
36     ...
37     def play_random_games(self):
38         for game in range(self.random_games):
39             self.env.reset()
40             for step in range(self.game_steps):
41                 if self.log:
42                     self.env.render()          #rendering the environment at each step (slow)
43                 action = self.env.action_space.sample()    #get a random action from the list of all possible actions
44                 observation, reward, done, info = self.env.step(action) #execute the action. returns lots of information
45                 if done:
46                     print("Game finished after {} steps".format(step+1))
47                     break
48
49     ...
50     To build our training set we must play some games!
```

Value

Value
ch = 9) completed out of 150 loss: 0.6165554
ch = 0) completed out of 150 loss: 0.60273606
ch = 1) completed out of 150 loss: 0.61319846
ch = 2) completed out of 150 loss: 0.5838175
ch = 3) completed out of 150 loss: 0.6168544
ch = 4) completed out of 150 loss: 0.61458105
Epoch 148 (batch = 5) completed out of 150 loss: 0.5808482
Epoch 148 (batch = 6) completed out of 150 loss: 0.60360914
Epoch 148 (batch = 7) completed out of 150 loss: 0.6458781
Epoch 148 (batch = 8) completed out of 150 loss: 0.62261486
Epoch 148 (batch = 9) completed out of 150 loss: 0.60551125
Epoch 149 (batch = 0) completed out of 150 loss: 0.6271517
Epoch 149 (batch = 1) completed out of 150 loss: 0.6277879
Epoch 149 (batch = 2) completed out of 150 loss: 0.60631627
Epoch 149 (batch = 3) completed out of 150 loss: 0.5965201
Epoch 149 (batch = 4) completed out of 150 loss: 0.6315039
Epoch 149 (batch = 5) completed out of 150 loss: 0.615431
Epoch 149 (batch = 6) completed out of 150 loss: 0.61090666

File explorer Help

IPython console History log

The CartPole-v1 Example

73

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

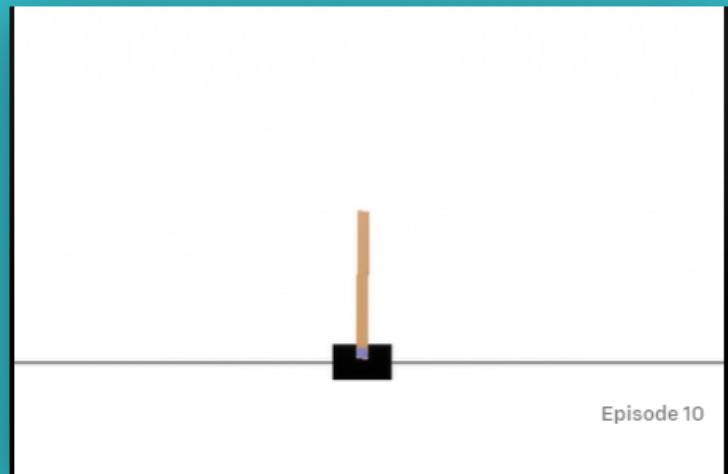
CartPole-v1

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The system is controlled by applying a force of +1 or -1 to the cart. The pendulum starts upright, and the goal is to prevent it from falling over. A reward of +1 is provided for every timestep that the pole remains upright. The episode ends when the pole is more than 15 degrees from vertical, or the cart moves more than 2.4 units from the center.

This environment corresponds to the version of the cart-pole problem described by Barto, Sutton, and Anderson [Barto83].

[Barto83] AG Barto, RS Sutton and CW Anderson, "Neuronlike Adaptive Elements That Can Solve Difficult Learning Control Problem", *IEEE Transactions on Systems, Man, and Cybernetics*, 1983.

↗ [VIEW SOURCE ON GITHUB](#)



RandomAgent on CartPole-v1

The CartPole-v1 Example

74

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

13 class CartPoleEnv(gym.Env):
14     """
15     Description:
16         A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The pendulum starts upright, and the ...
17
18     Source:
19         This environment corresponds to the version of the cart-pole problem described by Barto, Sutton, and Anderson
20
21     Observation:
22         Type: Box(4)
23             Num      Observation          Min          Max
24             0      Cart Position        -4.8          4.8
25             1      Cart Velocity        -Inf          Inf
26             2      Pole Angle           -24 deg       24 deg
27             3      Pole Velocity At Tip -Inf          Inf
28
29     Actions:
30         Type: Discrete(2)
31             Num      Action
32             0      Push cart to the left
33             1      Push cart to the right
34
35     Note: The amount the velocity that is reduced or increased is not fixed; it depends on the angle the pole is pointing. This is because ...
36
37     Reward:
38         Reward is 1 for every step taken, including the termination step
39
40     Starting State:
41         All observations are assigned a uniform random value in [-0.05..0.05]
42
43     Episode Termination:
44         Pole Angle is more than 12 degrees
45         Cart Position is more than 2.4 (center of the cart reaches the edge of the display)
46         Episode length is greater than 200
47         Solved Requirements
48             Considered solved when the average reward is greater than or equal to 195.0 over 100 consecutive trials.
49 """

```

(https://github.com/openai/gym/blob/master/gym/envs/classic_control/cartpole.py)

The CartPole-v1 Example

75

tf.data API

tf.keras API

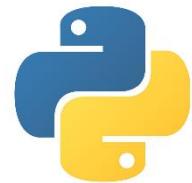
TensorBoard

OPENAI

Hands On

```
# =====
# Main Execution
#
# =====

playerAgent = PlayerAgent(500, 1000, 75, game='CartPole-v1', log=True)
#playerAgent.play_random_games()
playerAgent.play_the_game()
```



The CartPole-v1 Example

76

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

# =====
#
#     Player Agent
#
# =====
class PlayerAgent:

    ...
    Constructor - defining game variables
    ...

    def __init__(self, game_steps, nr_of_games_for_training, score_requirement, game='CartPole-v1', log=False):
        self.game_steps = game_steps                                #number of frames (aka steps) to play for each game
        self.nr_of_games_for_training = nr_of_games_for_training   #the number of games that will be played to build the training set
        self.score_requirement = score_requirement                  #the minimum score for a decent game (also for training purposes)
        self.random_games = 50                                      #initial random games to play
        self.log = log                                              #log info
        self.env = gym.make(game)                                    #make the gaming environment
        self.scores = []                                            #list with all scores when playing

    ...
    Let us test the environment and play some random games!
    ...

    def play_random_games(self):
        pass

    ...
    To build our training set we must play some games!
    Let us start by playing completely random games and save those where we achieve an interesting score.
    You should then use this training set to train the model!!
    ...

    def build_training_set(self):
        pass

    ...
    Let us play the game after been trained to do that!
    ...

    def play_the_game(self):
        pass

```

The CartPole-v1 Example

77

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```
# -----
# 
#     Player Agent
#
# -----
class PlayerAgent:
    ...
    Constructor - defining game variables
    ...
    def __init__(self, game_steps, nr_of_games_for_training, score_requirement, game='CartPole-v1', log=False):
        self.game_steps = game_steps                                #number of frames (aka steps) to play for each game
        self.nr_of_games_for_training = nr_of_games_for_training   #the number of games that will be played to build the training set
        self.score_requirement = score_requirement                 #the minimum score for a decent game (also for training purposes)
        self.random_games = 50                                      #initial random games to play
        self.log = log                                              #log info
        self.env = gym.make(game)                                    #make the gaming environment
        self.scores = []                                            #list with all scores when playing
```

- Start by calling `gym.make('SOME_GAME')` to prepare the game environment!
- Save it to a variable (let's call it `env`).

The CartPole-v1 Example

78

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```
'''  
Let us test the environment and play some random games!  
'''  
  
def play_random_games(self):  
    for game in range(self.random_games):  
        self.env.reset()  
        for step in range(self.game_steps):  
            if self.log:  
                self.env.render()  
            action = self.env.action_space.sample()  
            observation, reward, done, info = self.env.step(action)  
            if done:  
                print("Game finished after {} steps".format(step+1))  
                break  
        #let's play self.random_games games  
        #reset the environment to initial state. returns the first observation of the  
        #number of frames (aka steps) to play for each game  
  
        #rendering the environment at each step (slows things down)  
        #get a random action from the list of all possible actions  
        #execute the action. returns lots of information  
        #if game is finished. Time to reset the environment again
```

Let's play some random games (fool behaviour is expected!!)

The CartPole-v1 Example

79

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```
'''  
Let us test the environment and play some random games!  
'''  
  
def play_random_games(self):  
    for game in range(self.random_games):  
        self.env.reset()  
        for step in range(self.game_steps):  
            if self.log:  
                self.env.render()  
            action = self.env.action_space.sample()  
            observation, reward, done, info = self.env.step(action)  
            if do
```

C:\data\PythonWorkspace\mlfa_csc_tf2\Ficha4_OpenAI.py

— □ ×

t again



The CartPole-v1 Example

80

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

...
Let us test the environment and play some random games!
...

def play_random_games(self):
    for game in range(self.random_games):
        self.env.reset()                                #let's play self.random_games games
        for step in range(self.game_steps):             #reset the environment to initial state. returns the first observation of the
                                                        #number of frames (aka steps) to play for each game
            if self.log:                               #rendering the environment at each step (slows things down)
                self.env.render()                      #get a random action from the list of all possible actions
            action = self.env.action_space.sample()     #execute the action. returns lots of information
            observation, reward, done, info = self.env.step(action)  #if game is finished. Time to reset the environment again
            if done:
                print("Game finished after {} steps".format(step+1))
                break

```

- Call `env.reset()` to make the game environment ready to be played!
- Call `env.render()` to render the environment at each step (slows everything considerably)
- Call `env.action_space.sample()` to get a random action from all available actions (left | right)

The CartPole-v1 Example

81

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

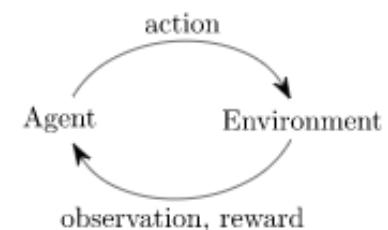
```

...
Let us test the environment and play some random games!
...

def play_random_games(self):
    for game in range(self.random_games):
        self.env.reset()                                #let's play self.random_games games
        for step in range(self.game_steps):             #reset the environment to initial state. returns the first observation of the
                                                       #number of frames (aka steps) to play for each game
            if self.log:                               #rendering the environment at each step (slows things down)
                self.env.render()
            action = self.env.action_space.sample()      #get a random action from the list of all possible actions
            observation, reward, done, info = self.env.step(action)  #execute the action. returns lots of information
            if done:                                  #if game is finished. Time to reset the environment again
                print("Game finished after {} steps".format(step+1))
                break

```

- Call `env.step(action)` to execute the *action* and get information about the consequence of having executed that same action!
 - **observation**: an environment-specific object representing your observation of the environment (check slide 74 to see what this environment returns)
 - **reward**: mount of reward achieved by the executed action
 - **done**: is time to reset the environment again?
 - **info**: useful for debugging!



The CartPole-v1 Example

82

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```
def build_training_set(self):
    training_set = []
    score_set = []

    for game in range(self.nr_of_games_for_training):
        cumulative_game_score = 0
        game_memory = []
        obs_prev = self.env.reset()
```

How can we train our agent so that it can play the game?

We must first collect some data!

The CartPole-v1 Example

83

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

def build_training_set(self):
    training_set = []
    score_set = []

    for game in range(self.nr_of_games_for_training):
        cumulative_game_score = 0
        game_memory = []
        obs_prev = self.env.reset()

        for step in range(self.game_steps):
            #env.render()
            action = self.env.action_space.sample()
            obs_next, reward, done, info = self.env.step(action)
            game_memory.append([action, obs_prev])
            cumulative_game_score += reward
            obs_prev = obs_next
            if done:
                #print("Game finished after {} steps".format(step+1))    #for debug purposes
                break

```

#best not to render the game!
#some random action
#execute it and gather the results
#store the execute action and the corresponding observation
#store cumulative reward per game

How can we train our agent so that it can play the game?

We must first collect some data!

The CartPole-v1 Example

84

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

def build_training_set(self):
    training_set = []
    score_set = []

    for game in range(self.nr_of_games_for_training):
        cumulative_game_score = 0
        game_memory = []
        obs_prev = self.env.reset()

        for step in range(self.game_steps):
            #env.render()
            action = self.env.action_space.sample()
            obs_next, reward, done, info = self.env.step(action)
            game_memory.append([action, obs_prev])
            cumulative_game_score += reward
            obs_prev = obs_next
            if done:
                #print("Game finished after {} steps".format(step+1))    #for debug purposes
                break

        #the game is finished. Was it a decent game?
        if cumulative_game_score > self.score_requirement:
            score_set.append(cumulative_game_score)
            for play in game_memory:
                if play[0] == 0:
                    one_hot_action = [1, 0]
                elif play[0] == 1:
                    one_hot_action = [0, 1]

            training_set.append([one_hot_action, play[1]])

```

#best not to render the game!
#some random action
#execute it and gather the results
#store the execute action and the corresponding observation
#store cumulative reward per game

How can we train our agent so that it can play the game?

We must first collect some data!

The CartPole-v1 Example

85

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

def build_training_set(self):
    training_set = []
    score_set = []

    for game in range(self.nr_of_games_for_training):
        cumulative_game_score = 0
        game_memory = []
        obs_prev = self.env.reset()

        for step in range(self.game_steps):
            #env.render()
            action = self.env.action_space.sample()
            obs_next, reward, done, info = self.env.step(action)
            game_memory.append([action, obs_prev])
            cumulative_game_score += reward
            obs_prev = obs_next
            if done:
                #print("Game finished after {} steps".format(step+1))    #for debug purposes
                break

        #the game is finished. Was it a decent game?
        if cumulative_game_score > self.score_requirement:
            score_set.append(cumulative_game_score)
            for play in game_memory:
                if play[0] == 0:
                    one_hot_action = [1, 0]
                elif play[0] == 1:
                    one_hot_action = [0, 1]

            training_set.append([one_hot_action, play[1]])

    #print some stats
    if score_set:
        print('Average score:', mean(score_set))
        print('Number of stored games per score', Counter(score_set))

    return training_set

```

#best not to render the game!
#some random action
#execute it and gather the results
#store the execute action and the corresponding observation
#store cumulative reward per game

How can we train our agent so that it can play the game?

We must first collect some data!

The CartPole-v1 Example

86

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```

...
Let us play the game after been trained to do that!
...

def play_the_game(self):
    #build training set
    training_set = self.build_training_set()

    #using the model
    mlp = MLP()
    mlp.build()
    mlp.fit(training_set)

    for game in range(self.nr_of_games_for_training):
        score = 0
        obs_prev = self.env.reset()
        done = False

        while not done:
            if self.log:
                self.env.render()
            action = mlp.predict(obs_prev)
            obs_next, reward, done, _ = self.env.step(action)
            score += reward
            obs_prev = obs_next

        self.scores.append(score)
        print('Current score: ' + str(score) + '; Average score:', sum(self.scores)/len(self.scores))

```

To play the game we (1) **collect the training set** (or persist first and load later), (2) **build our MLP**, (3) **train** it with the training set, and (4) we start **playing the game** using the action returned by our model when it receives the current observation of the environment!

The CartPole-v1 Example

87

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On



Much better!! But it can be (substantially) improved!!!

How?

The CartPole-v1 Example

88

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

```
# =====
# Multi-Layer Perceptron
#
# =====
class MLP:

    ...
    Constructor - defining game variables
    ...

    def __init__(self, epochs=1, batch_size=32, output_neurons=2):
        self.epochs = epochs
        self.batch_size = batch_size
        self.output_neurons = output_neurons

    def prepare_data(self, training_set):
        #reshape X to list of inputs
        X = np.array([i[1] for i in training_set], dtype="float32").reshape(-1, len(training_set[0][1]))
        #y are the actions
        y = np.array([i[0] for i in training_set])
        return X, y

    def build(self):
        pass

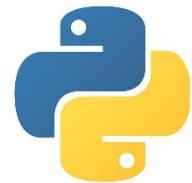
    def fit(self, training_set):
        pass

    def predict(self, obs):
        pass
```

Now... You just need to build a MLP!!

OpenAI Gym

The CartPole-v1 Example



89

tf.data API

tf.keras API

TensorBoard

OPENAI

Hands On

Spyder (Python 3.6)

File Edit Search Source Run Debug Consoles Projects Tools View Help

Editor - C:/data/PythonWorkspace/mlfa_csc/TP1/TP1_example.py

TP1_example.py

```
15 # -----
16 #
17 #           Player Agent
18 #
19 # -----
20 class PlayerAgent:
21     ...
22     Constructor - defining game variables
23     ...
24     def __init__(self, game_steps, nr_of_games_for_training):
25         self.game_steps = game_steps
26         self.nr_of_games_for_training = nr_of_games_for_training
27         self.score_requirement = score_requirement
28         self.random_games = 50
29         self.log = log
30         self.env = gym.make(game)
31         self.scores = []
32     ...
33     ...
34     Let us test the environment and play some games!
35     ...
36     def play_some_random_games(self):
37         for game in range(self.random_games):
38             self.env.reset()
39             for step in range(self.game_steps):
40                 if self.log:
41                     self.env.render()          #rendering the environment at each step (slow)
42                 action = self.env.action_space.sample()    #get a random action from the list of all possible actions
43                 observation, reward, done, info = self.env.step(action) #execute the action. returns lots of information
44                 if done:
45                     print("Game finished after {} steps".format(step+1))
46                     break
47             ...
48     ...
49     To build our training set we must play some games!
```

C:/data/PythonWorkspace/mlfa_csc/TP1/TP1_example.py

Value

File explorer Help

ch = 9) completed out of 150 loss: 0.6165554
ch = 0) completed out of 150 loss: 0.60273606
ch = 1) completed out of 150 loss: 0.61319846
ch = 2) completed out of 150 loss: 0.5838175
ch = 3) completed out of 150 loss: 0.6168544
ch = 4) completed out of 150 loss: 0.61458105
Epoch 148 (batch = 5) completed out of 150 loss: 0.5808482
Epoch 148 (batch = 6) completed out of 150 loss: 0.60360914
Epoch 148 (batch = 7) completed out of 150 loss: 0.6458781
Epoch 148 (batch = 8) completed out of 150 loss: 0.62261486
Epoch 148 (batch = 9) completed out of 150 loss: 0.60551125
Epoch 149 (batch = 0) completed out of 150 loss: 0.6271517
Epoch 149 (batch = 1) completed out of 150 loss: 0.6277879
Epoch 149 (batch = 2) completed out of 150 loss: 0.60631627
Epoch 149 (batch = 3) completed out of 150 loss: 0.5965201
Epoch 149 (batch = 4) completed out of 150 loss: 0.6315039
Epoch 149 (batch = 5) completed out of 150 loss: 0.615431
Epoch 149 (batch = 6) completed out of 150 loss: 0.6109066

IPython console History log

Back to Spyder

Glossary



90

tf.data API

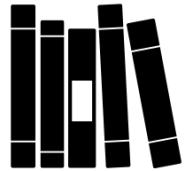
tf.keras API

TensorBoard

OpenAI

HANDS ON

- **Callbacks**: an object passed to a model to customize and extend its behavior during training
- **Checkpoints**: a file that captures the values of all `tf.Variable` objects used by a model allowing the ability to stop the training at any point, for any reason, and resume it later as if nothing happened
- **Keras**: a high-level neural networks API, written in Python and capable of running on top of TensorFlow. `tf.Keras` is TensorFlow's official high-level API
- **OpenAI Gym**: a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball
- **Tensorboard**: a set of visualization tools to make it easier to understand, debug and optimize TensorFlow programs
- **`tf.Data`**: a TensorFlow API to handle large amounts of data, read from different data formats, and perform complex transformations



Resources

91

tf.data API

tf.keras API

TensorBoard

OpenAI

HANDS ON

- Official Documentation
 - https://www.tensorflow.org/api_docs
 - <https://www.tensorflow.org/guide/data>
 - https://www.tensorflow.org/tutorials/load_data
 - https://www.tensorflow.org/guide/keras/custom_callback
 - https://www.tensorflow.org/guide/saved_model
 - <https://gym.openai.com/docs/>
 - <https://github.com/openai/gym>
 - ...

Hands On



92

tf.data API

tf.keras API

TensorBoard

OpenAI

HANDS ON

● Spyder (Python 3.6)

File Edit Search Source Run Debug Consoles Projects Tools View Help

Editor - C:\data\PythonWorkspace\dev\meanshift_algorithm.py

deep-network.py meanshift_algorithm.py

```
37
38 class Mean_Shift:
39     def __init__(self, radius=None, radius_normalize_step = 150):
40         self.radius = radius
41         self.radius_normalize_step = radius_normalize_step
42
43     def fit(self, data):
44
45         if self.radius == None:
46             all_data_centroid = np.average(data, axis=0)
47             all_data_norm = np.linalg.norm(all_data_centroid)
48             self.radius = all_data_norm/self.radius_normalize_step
49
50         centroids = {}
51
52         #initialize centroids
53         for i in range(len(data)):
54             centroids[i] = data[i]
55
56         weights = [i for i in range(self.radius_normalize_step)]
57
58         while True:
59             new_centroids = []
60             for i in centroids:
61                 in_range = []
62                 centroid = centroids[i]
63
64                 for featureset in data:
65                     distance = np.linalg.norm(featureset-centroid)
66                     if distance == 0:
67                         distance = 0.0000000001
68                     weight_index = int(distance/self.radius)
69                     if weight_index > self.radius_normalize_step-1:
70                         weight_index = self.radius_normalize_step-1
71                     to_add = (weights[weight_index]**2)[featureset]
72                     in_range += to_add
73
74             new_centroid = np.average(in_range, axis=0)
```

Variable explorer

Name	Type	Size	Value
batch_size	int	1	100
mnist	contrib.learn.python.learn.datasets.base.Datasets	3	Datasets object of...
n_classes	int	1	10
n_nodes_hl1	int	1	500
n_nodes_hl2	int	1	500
n_nodes_hl3	int	1	500

Variable explorer File explorer Help

IPython console

In [2]:

```
Epoch 0 completed out of 10 loss: 1666037.4677734375
Epoch 1 completed out of 10 loss: 377809.3128890991
Epoch 2 completed out of 10 loss: 201302.4857263565
Epoch 3 completed out of 10 loss: 119427.91378033161
Epoch 4 completed out of 10 loss: 72651.25679710507
Epoch 5 completed out of 10 loss: 45327.621502393486
Epoch 6 completed out of 10 loss: 31955.17812934518
Epoch 7 completed out of 10 loss: 23664.356108333137
Epoch 8 completed out of 10 loss: 18248.740643078025
Epoch 9 completed out of 10 loss: 19962.00085876091
Accuracy: 0.9511
```

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

HANDS ON