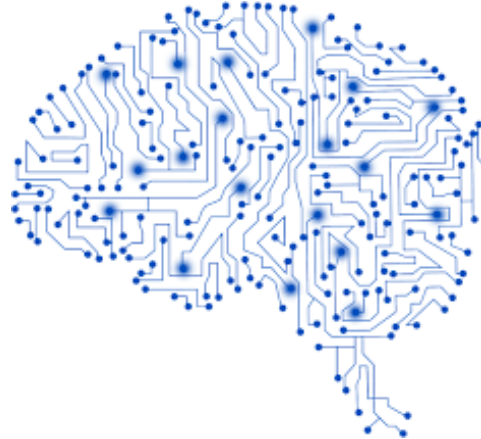




University of Minho
School of Engineering



LSTM networks with TensorFlow™

Connective Systems and Classifiers

Perfil ML:FA @ MiEI/4º ano - 2º Semestre

Bruno Fernandes, Victor Alves

12/03/2020

Contents

2

LSTMs

Seq. Prediction

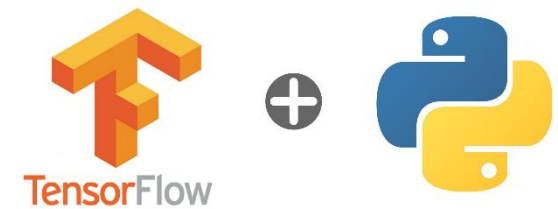
Preparing Data

TF and LSTMs

Hands On

- Long Short-Term Memory Networks
- Sequence Prediction Problems
- Preparing data for LSTMs
- Developing LSTM networks
- The Echo Sequence Prediction Problem
- Hands On

LSTM Networks



3

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

Hands On

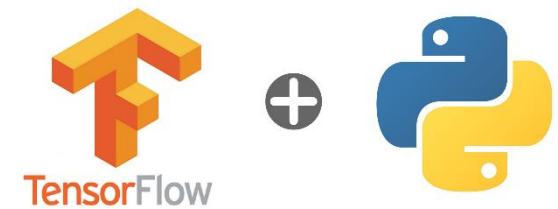
MLPs are a good starting point for modeling sequence prediction problems ... But we now have better options!

The **Long Short-Term Memory**, or **LSTM**, network is a type of **Recurrent Neural Network**, or RNN, specially designed for sequence problems. The promise of RNNs, and LSTMs in particular, is that the **temporal dependence** and **contextual information** in the input data can be **learned**!

Q.: When were **LSTMs introduced**?

A.: In 1997, in the work of Hochreiter & Schmidhuber, but many more contributed to the modern LSTM. Check the original paper [here](#)!

LSTM Networks



4

LSTMs

Seq. Prediction

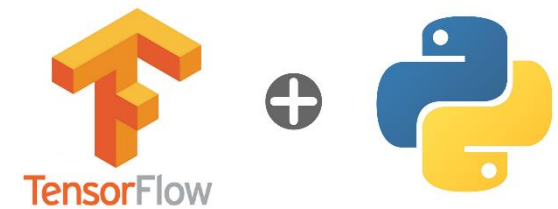
Preparing Data

TF and LSTMs

Hands On

- The computational unit of the LSTM network is called the **memory cell** (memory block, neuron or just cell are also used)!
- LSTM neurons are comprised of **weights** and **gates**
- The key to the LSTM neurons (memory cells) are the gates. These too are **weighted functions** that further govern the **information flow** (internal state) in the cell.

LSTM Networks



5

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

Hands On

- A neuron has three gates:
 - **Forget Gate** that decides what information to **discard** from the **internal state**
 - **Input Gate** that decides which values from the input to **add** to the **internal state**
 - **Output Gate** that decides what to **output** based on the **input and internal state**
- The **forget** and **input** gates are used to update the internal state of the neuron
- The **output** gate is a final limiter on what the cell actually outputs.
- It is these gates and the consistent data flow that keep each cell stable (avoiding the **vanishing** and **exploding gradients problem**)

A memory cell



6

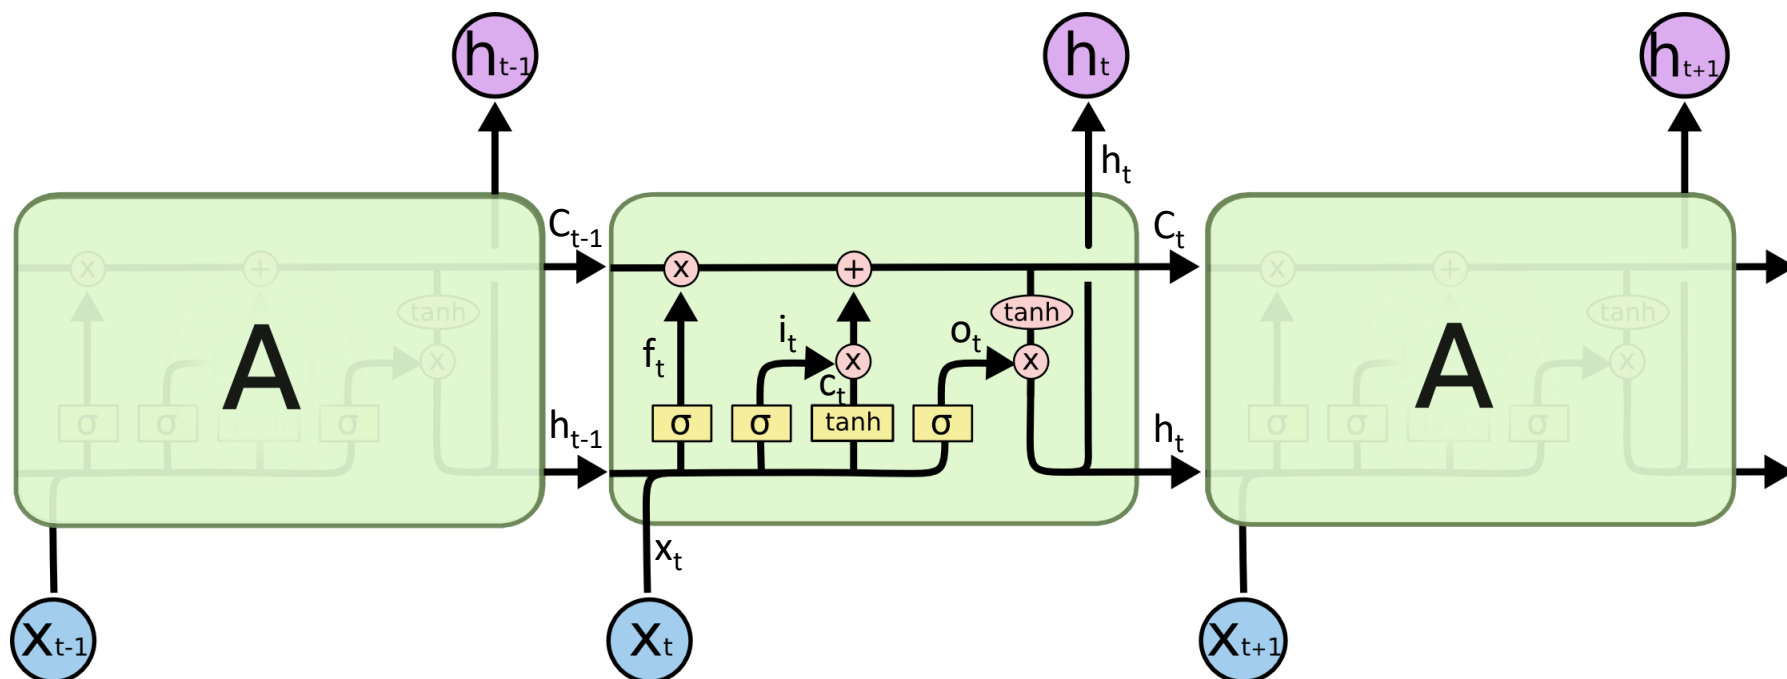
LSTMs

Seq. Prediction

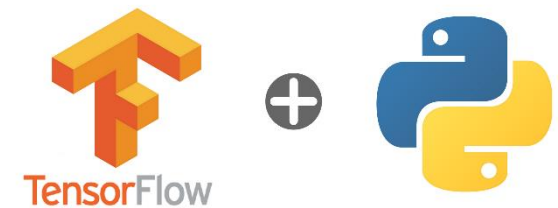
Preparing Data

TF and LSTMs

Hands On



Application of LSTMs



7

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

Hands On

Automatic Image Caption Generation

A person riding a motorcycle on a dirt road.



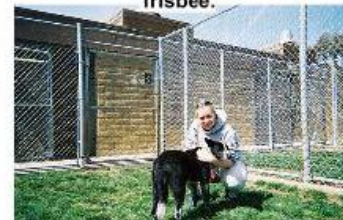
Two dogs play in the grass.



A skateboarder does a trick on a ramp.



A dog is jumping to catch a frisbee.



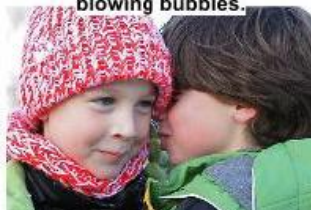
A group of young people playing a game of frisbee.



Two hockey players are fighting over the puck.



A little girl in a pink hat is blowing bubbles.



A refrigerator filled with lots of food and drinks.



A herd of elephants walking across a dry grass field.



A close up of a cat laying on a couch.



A red motorcycle parked on the side of the road.



A yellow school bus parked in a parking lot.



Describes without errors

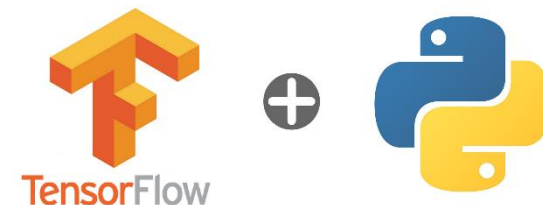
Describes with minor errors

Somewhat related to the image

Unrelated to the image

(Show and Tell: A Neural Image Caption Generator, 2014)

Application of LSTMs



8

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

Hands On

Automatic Translation of Text

Type	Sentence
Our model	Ulrich UNK , membre du conseil d' administration du constructeur automobile Audi , affirme qu' il s' agit d' une pratique courante depuis des années pour que les téléphones portables puissent être collectés avant les réunions du conseil d' administration afin qu' ils ne soient pas utilisés comme appareils d' écoute à distance .
Truth	Ulrich Hackenberg , membre du conseil d' administration du constructeur automobile Audi , déclare que la collecte des téléphones portables avant les réunions du conseil , afin qu' ils ne puissent pas être utilisés comme appareils d' écoute à distance , est une pratique courante depuis des années .
Our model	“ Les téléphones cellulaires , qui sont vraiment une question , non seulement parce qu' ils pourraient potentiellement causer des interférences avec les appareils de navigation , mais nous savons , selon la FCC , qu' ils pourraient interférer avec les tours de téléphone cellulaire lorsqu' ils sont dans l' air ” , dit UNK .
Truth	“ Les téléphones portables sont véritablement un problème , non seulement parce qu' ils pourraient éventuellement créer des interférences avec les instruments de navigation , mais parce que nous savons , d' après la FCC , qu' ils pourraient perturber les antennes-relais de téléphonie mobile s' ils sont utilisés à bord ” , a déclaré Rosenker .
Our model	Avec la crémation , il y a un “ sentiment de violence contre le corps d' un être cher ” , qui sera “ réduit à une pile de cendres ” en très peu de temps au lieu d' un processus de décomposition “ qui accompagnera les étapes du deuil ” .
Truth	Il y a , avec la crémation , “ une violence faite au corps aimé ” , qui va être “ réduit à un tas de cendres ” en très peu de temps , et non après un processus de décomposition , qui “ accompagnerait les phases du deuil ” .

Application of LSTMs



9

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

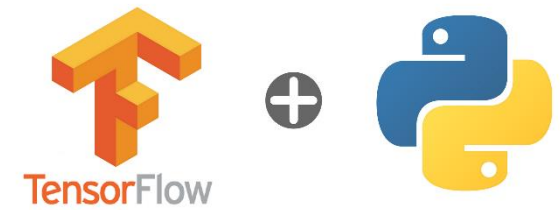
Hands On

Automatic Handwriting Generation

from his travels it might have been
from his travels it might have been
from his travels it might have been

from his travels it might have been
from his travels it might have been
from his travels it might have been

Sequence Prediction



10

LSTMs

SEQ. PREDICTION

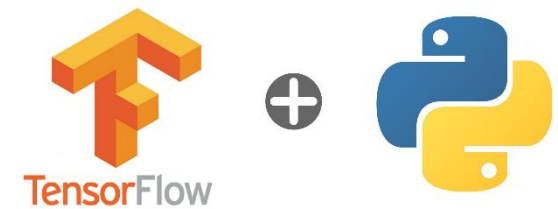
Preparing Data

TF and LSTMs

Hands On

- **Sequence prediction** is different to other types of supervised learning problems
- The sequence imposes an **explicit order on the observations** that must be **preserved** when training models and making predictions
- **Order is important** and it must be respected in the formulation of prediction problems that use the sequence data as input!

Sequence Prediction



11

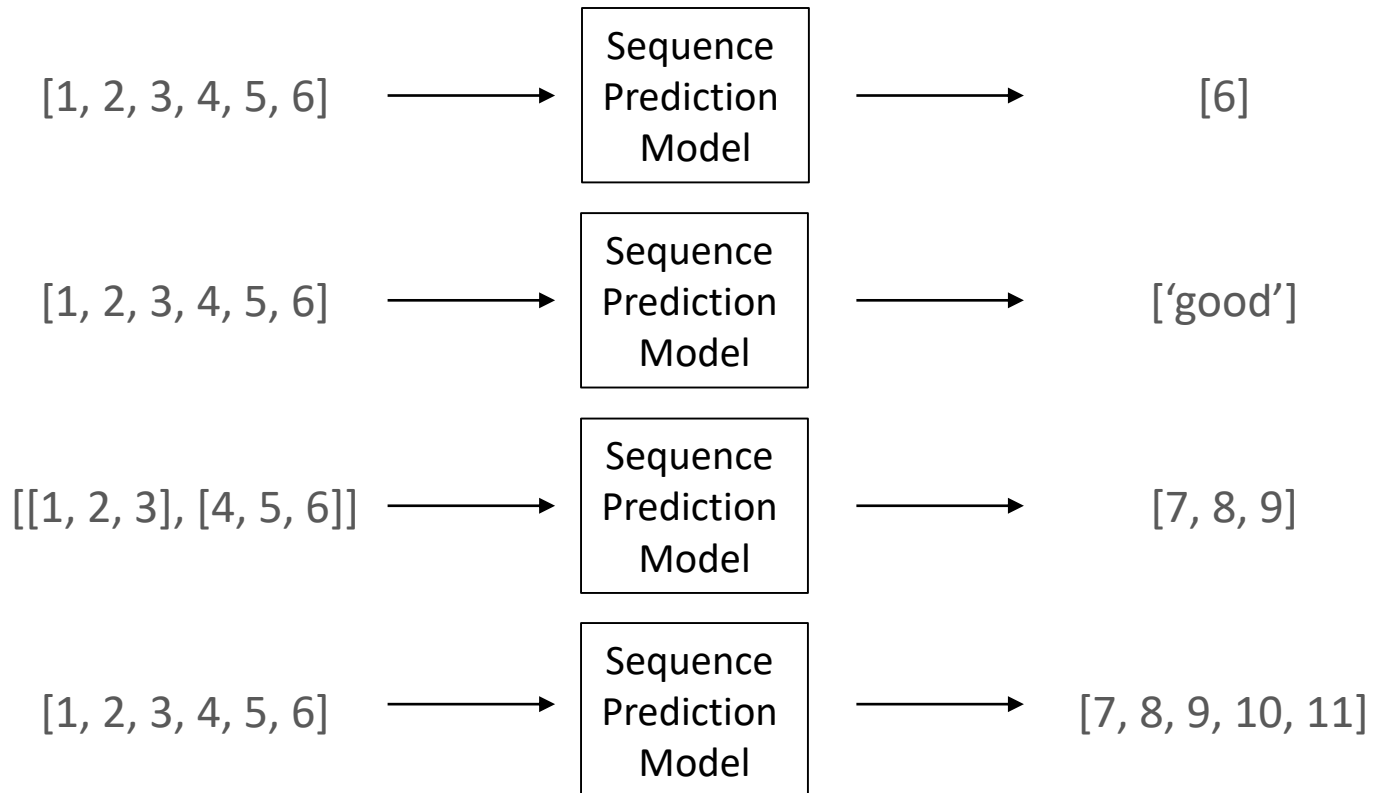
LSTMs

SEQ. PREDICTION

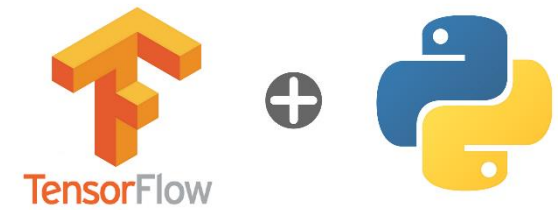
Preparing Data

TF and LSTMs

Hands On



Sequence Prediction



12

LSTMs

SEQ. PREDICTION

Preparing Data

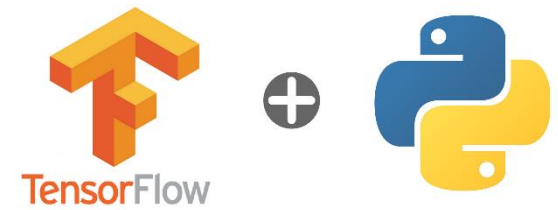
TF and LSTMs

Hands On

Single-step vs Multi-step

- Single-step consists in forecasting just **the next value** of the sequence (**next-element prediction**)
 - Given the observed temperature over the last 7 days forecast the temperature at timestep 8
- Multi-step consists in forecasting **multiple timesteps** beyond the input sequences (**sequence-to-sequence prediction**)
 - Given the observed temperature over the last 7 days forecast the temperature at timestep 8, 9 and 10
 - Accurate straightforward multi-step forecasting remains a challenge to our days

Sequence Prediction



13

LSTMs

SEQ. PREDICTION

Preparing Data

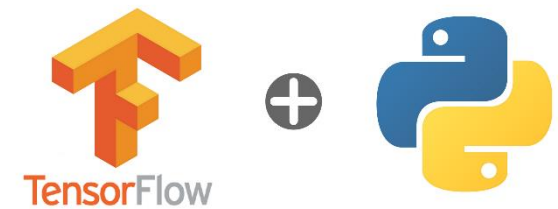
TF and LSTMs

Hands On

Univariate vs Multivariate

- Univariate, or single input, is where you use **a single feature** as input over multiple input timesteps
 - Given the observed pollution in previous timesteps predict the expected pollution value
- Multivariate, or multiple input, is where you use **multiple features** as input over multiple input timesteps
 - Given the observed pollution and weather in previous timesteps predict the expected pollution value

Preparing Data for LSTMs



14

LSTMs

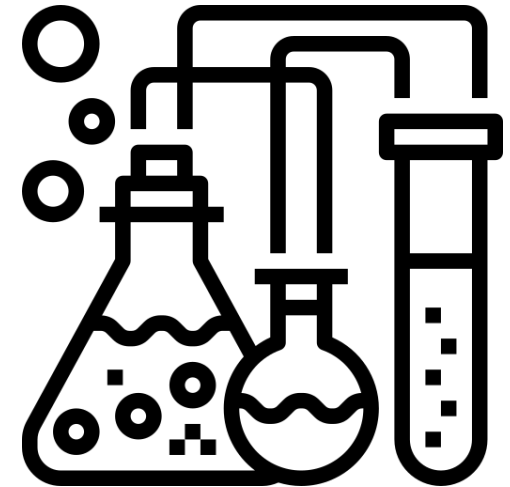
Seq. Prediction

PREPARING DATA

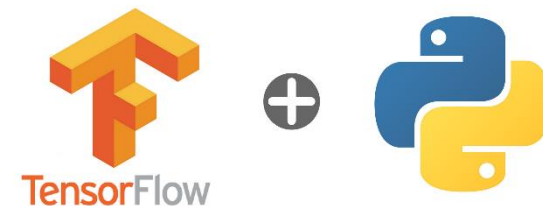
TF and LSTMs

Hands On

1. Preparing **Numeric** Data
2. Preparing **Categorical** Data
3. Preparing **Sequences** with **Varied Lengths**
4. Preparing **Sequence Prediction** as **Supervised Learning**



Preparing Numeric Data



15

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

When a network is fit on **unscaled data**, it is possible for large inputs to **slow down** the learning and convergence of the network! It may **prevent the network from learning** our problem!

There are two types of scaling that we may want to consider (use scikit-learn):

- **Normalization**
 - rescaling data so that all values fall within the range of 0 and 1
- **Standardization**
 - rescaling the distribution of values so that the mean of observed values is 0 and the standard deviation is 1
 - assumes observations fit a Gaussian distribution with a well behaved mean and standard deviation - which may not always be the case

Normalization



16

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Normalization can be achieved using the **scikit-learn** object **MinMaxScaler**

```
from pandas import Series
from sklearn.preprocessing import MinMaxScaler

#define series
series = Series([0.0, 10.0, 20.0, 30.0, 40.0, 50.0, 60.0, 70.0, 80.0,
90.0, 100.0])

#prepare data for normalization
values = series.values.reshape((len(series.values), 1))

#fit and transform or fit_transform
scaler = MinMaxScaler(feature_range=(0, 1))
normalized = scaler.fit_transform(values)
print(normalized)

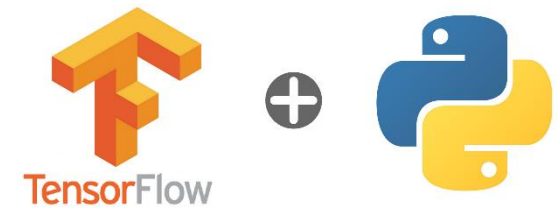
#inverse transform is also a thing
inversed = scaler.inverse_transform(normalized)
print(inversed)
```



```
[[0. ]
 [0.1]
 [0.2]
 ...
 [0.9]
 [1. ]]

[[ 0.]
 [10.]
 [20.]
 ...
 [90.]
 [100.]]
```

Standardization



17

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Standardization can be achieved using the **scikit-learn** object **StandardScaler**

```
from pandas import Series
from sklearn.preprocessing import StandardScaler

#define series
series = Series([4.1, 1.0, 2.6, 7.9, 5.5, 9.0, 8.8, 3.0, 6.3])

#prepare data for standardization
values = series.values.reshape((len(series.values), 1))

#fit and transform or fit_transform
scaler = StandardScaler()
standardized = scaler.fit_transform(values)
print(standardized)

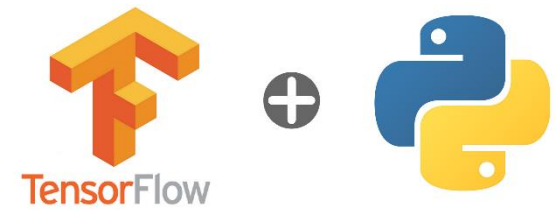
#inverse transform is still a thing
inversed = scaler.inverse_transform(standardized)
print(inversed)
```



```
[[ -0.46286604]
 [ -1.60569456]
 [ -1.01584758]
 ...
 [ -0.86838584]
 [  0.34817357]]

[[4.1]
 [1. ]
 [2.6]
 ...
 [3. ]
 [6.3]]
```

Preparing Categorical Data



18

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

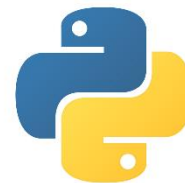
Hands On

Categorical data, often called nominal data, are variables that contain **label values** rather than numeric values

There are two types of treatment we may want to consider:

- **Label** (or integer) **Encoding**
 - integer values have a natural ordered relationship between each other and ML algorithms may be able to understand and harness this relationship
- **One Hot Encoding**
 - categorical variables where no such ordinal relationship exists

Preparing Categorical Data



19

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

It can be achieved using the **scikit-learn** objects **LabelEncoder** and **OneHotEncoder**

```
from numpy import array, argmax
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

#define array
values = array(['cold', 'warm', 'cold', 'hot', 'warm', 'warm', 'hot'])

#label encoding
label_encoder = LabelEncoder()
label_encoded = label_encoder.fit_transform(values)
print(label_encoded)

#one hot encoding
onehot_encoder = OneHotEncoder(sparse=False, categories='auto')
label_encoded = label_encoded.reshape(len(label_encoded), 1)
onehot = onehot_encoder.fit_transform(label_encoded)
print(onehot)

#inverse transform keeps being a thing
inverted = label_encoder.inverse_transform([argmax(onehot[0, :])])
print(inverted)
```



[0 2 0 1 2 2 1]

[[1. 0. 0.]

[0. 0. 1.]

[1. 0. 0.]

[0. 1. 0.]

[0. 0. 1.]

[0. 0. 1.]

[0. 1. 0.]

['cold']

Sequence with Varied Lengths



20

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Deep learning libraries assume a **vectorized representation of data**. In the case of variable length sequence prediction problems, sequences **should have the same length**

There are some types of treatment we may want to consider:

- **Padding** (adding a number, usually 0, to the beginning or the end of a sequence)
 - Pre-Sequence Padding
 - Post-Sequence Padding
- **Truncation** (trim a sequence to a desired length from its beginning or end)
 - Pre-Sequence Truncation
 - Post-Sequence Truncation

Padding Sequences



21

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Use `tf.keras.pad_sequences` object

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
#define sequences
```

```
sequences = [
```

```
[1, 2, 3, 4],
```

```
[1, 2, 3],
```

```
[1]
```

```
]
```

```
#pre-sequence padding
```

```
pre_padded = pad_sequences(sequences)
```

```
print(pre_padded)
```

```
#post-sequence padding
```

```
post_padded = pad_sequences(sequences, padding='post')
```

```
print(post_padded)
```



```
[[1 2 3 4]
```

```
[0 1 2 3]
```

```
[0 0 0 1]]
```

```
[[1 2 3 4]
```

```
[1 2 3 0]
```

```
[1 0 0 0]]
```

Truncating Sequences



22

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Use `tf.keras.pad_sequences` object and define the desired length on the `maxlen` argument

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

```
#define sequences
```

```
seq = [  
    [1, 2, 3, 4],  
    [1, 2, 3],  
    [1]  
]
```

```
#pre-sequence truncation
```

```
pre_truncated = pad_sequences(seq, maxlen=2)
```

```
print(pre_truncated)
```

```
#post-sequence truncation
```

```
post_truncated = pad_sequences(seq, maxlen=2, truncating='post')
```

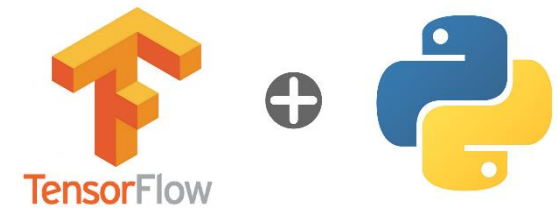
```
print(post_truncated)
```



```
[[3 4]  
 [2 3]  
 [0 1]]
```

```
[[1 2]  
 [1 2]  
 [0 1]]
```


Sequence Prediction as Supervised Learning



23

LSTMs

Seq. Prediction

PREPARING DATA

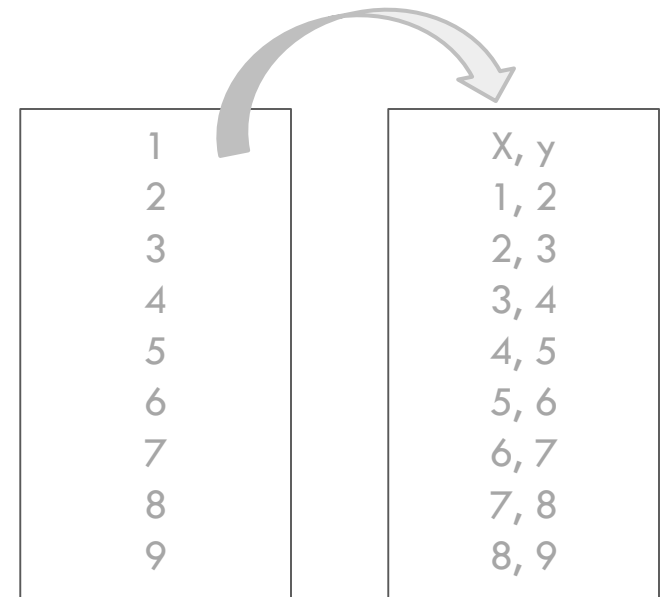
TF and LSTMs

Hands On

Sequence prediction problems must be re-framed as Supervised Learning problems!

Data must be transformed from a sequence to pairs of input/output. As we already know, a supervised learning problem allows models to learn how to **predict output patterns from input data**.

Pandas shift() function may help transform time series data into a supervised learning problem! Given a DataFrame, the `shift()` function can be used to **create copies of columns** that are **pushed forward** (rows of NaN values added to the front) or **pulled back** (rows of NaN values added to the end)



Sequence Prediction as Supervised Learning



24

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

```
from pandas import DataFrame
```

```
#define a mock timeseries as a column and as a sequence of 10 nrs
```

```
df = DataFrame()
```

```
df['t'] = [x for x in range(10)]
```

```
print(df)
```



```
t  
0 0  
1 1  
2 2  
...  
9 9
```

Sequence Prediction as Supervised Learning



25

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

```
from pandas import DataFrame
```

```
#define a mock timeseries as a column and as a sequence of 10 nrs
```

```
df = DataFrame()
```

```
df['t'] = [x for x in range(10)]
```

```
print(df)
```

```
#shift all obs down by 1 timestep by inserting a new row at the top
```

```
#the new row has no data so it will be represented as NaN
```

```
df['t-1'] = df['t'].shift(1)
```

```
print(df)
```



	t	
0	0	
1	1	
2	2	
...		
9	9	

	t	t-1
0	0	NaN
1	1	0.0
2	2	1.0
...		
9	9	8.0

Sequence Prediction as Supervised Learning



26

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

```
from pandas import DataFrame
```

```
#define a mock timeseries as a column and as a sequence of 10 nrs
```

```
df = DataFrame()
```

```
df['t'] = [x for x in range(10)]
```

```
print(df)
```

```
#shift all obs down by 1 timestep by inserting a new row at the top
```

```
#the new row has no data so it will be represented as NaN
```

```
df['t-1'] = df['t'].shift(1)
```

```
print(df)
```

```
#shift backward using a negative int value
```

```
#this pulls the observations up by inserting a new row at the end
```

```
df['t+1'] = df['t'].shift(-1)
```

```
print(df)
```



	t
0	0
1	1
2	2
...	
9	9

	t	t-1
0	0	NaN
1	1	0.0
2	2	1.0
...		
9	9	8.0

	t	t-1	t+1
0	0	NaN	1.0
1	1	0.0	2.0
2	2	1.0	3.0
...			
9	9	8.0	NaN

Sequence Prediction as Supervised Learning



27

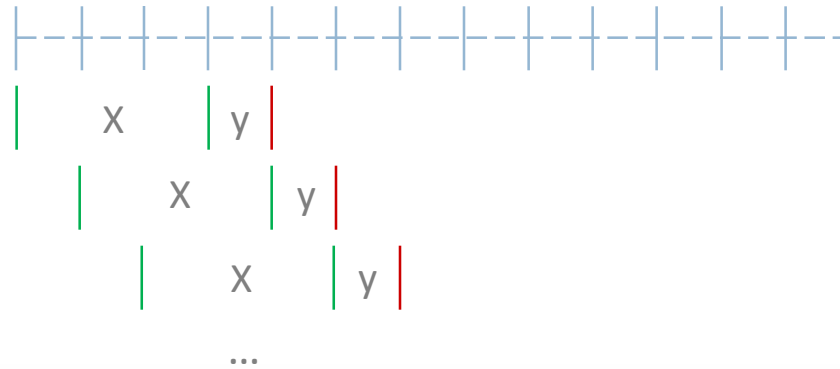
LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On



Algorithm 2: From an unsupervised to a supervised problem.

Input: dataset, timesteps, multi_steps

Output: input data, label data

$X, y = \text{list}(), \text{list}()$

while $i \in \text{range}(\text{len}(\text{dataset}) - (\text{timesteps} + \text{multi_steps}))$ **do**

 input_index = $i + \text{timesteps}$

 label_index = input_index + multi_steps

$X.\text{append}(\text{dataset}[i:\text{input_index}, :])$

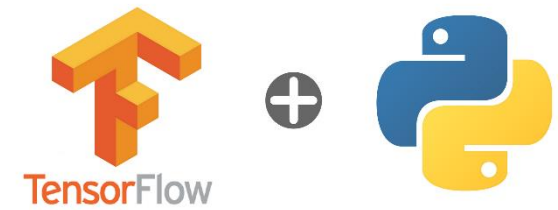
$y.\text{append}(\text{dataset}[\text{input_index}:\text{label_index}, \text{speed_diff}])$

end

return X, y

Sequence Prediction

Missing Timesteps



28

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Missing timesteps may be problematic for RNNs! There are several ways of handling those:

- **Masking** with default values
- **Interpolation**
- **Model-based computation**

Sequence Prediction

Missing Timesteps



29

LSTMs

Seq. Prediction

PREPARING DATA

TF and LSTMs

Hands On

Algorithm 1: Computation of missing timesteps.

```
dataset.resample('20min')
dataset['temperature', 'precipitation'].interpolate(method='linear',
    limit_direction='forward', limit='6h')
initialize consequent_missing_obs
foreach  $i, row \in enumerate(dataset, 0)$  do
    if row['speed_diff'] is NaN and consequent_missing_obs < 10h then
        increment consequent_missing_obs
        if  $i - (3 \times one\_week) > 0$  then
            value_1week_before = dataset[i - one_week]
            value_2weeks_before = dataset[i - one_week  $\times$  2]
            value_3weeks_before = dataset[i - one_week  $\times$  3]
            row['speed_diff'] = mean(value_1week_before,
                value_2weeks_before, value_3weeks_before)
        else
            value_1week_after = dataset[i + one_week]
            value_2weeks_after = dataset[i + one_week  $\times$  2]
            value_3weeks_after = dataset[i + one_week  $\times$  3]
            row['speed_diff'] = mean(value_1week_after, value_2weeks_after,
                value_3weeks_after)
        end
    else if row['speed_diff'] is not NaN then
        reset consequent_missing_obs
    end
end
```


Developing LSTM networks



30

LSTMs

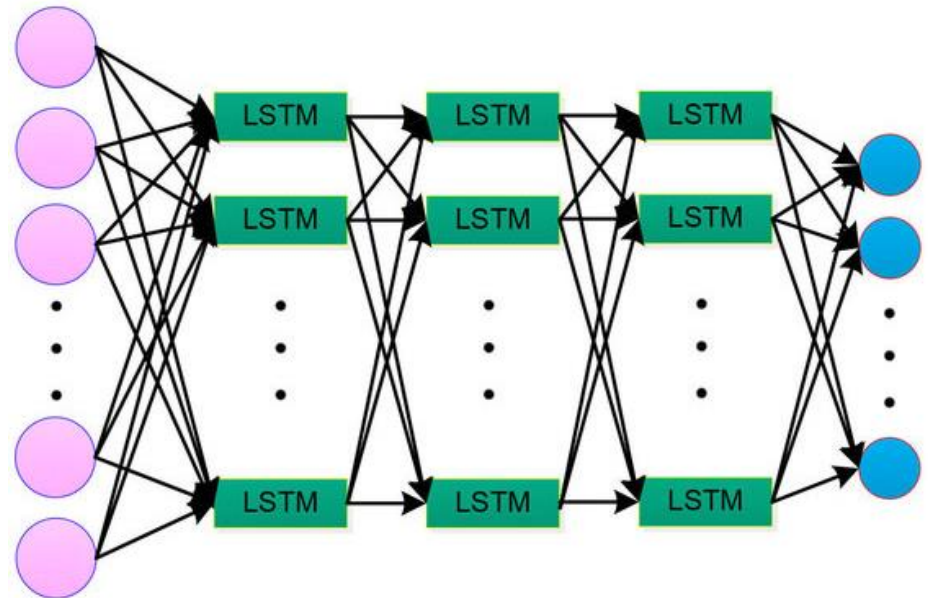
Seq. Prediction

Preparing Data

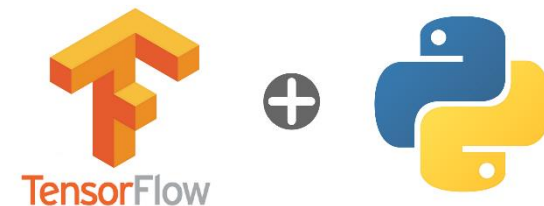
TF AND LSTMs

Hands On

1. The **Shape of Inputs**
2. **Shuffling Samples** (Not)
3. The **Importance of Batches** in LSTMs
4. LSTM's **State Management**
5. The **echo prediction problem**



The Shape of Inputs



31

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

The first step is to **define the model**! It is in the first LSTM layer **where we define the shape of the inputs** the model is expecting! **Input must be a 3D vector** comprising:

- **Samples**
 - One sequence is a sample! A batch is comprised of one, or more, samples
- **Timesteps**
 - Past observations of a feature. One timestep is one point of observation in the sample.
- **Features**
 - These are columns in your data. One feature is one observation at a timestep.

Use numpy **reshape()** function!

```
from numpy import array
```

```
data = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0])
```

```
print(data.shape)
```

```
data = data.reshape((1, 10, 1))
```

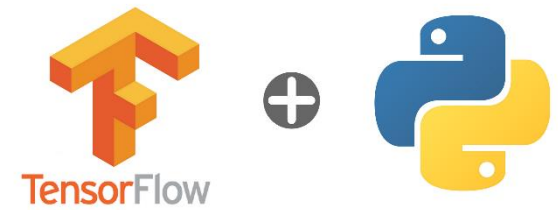
```
print(data.shape)
```



(10,)

(1, 10, 1)

The Shape of Inputs



32

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

The **first LSTM layer** in the network must **define the shape of the inputs it is expecting!** We need to define a tuple containing the **number of timesteps** and the **number of features!**

For example, if we had 10 timesteps and 1 feature (univariate sequence) per sample, the model would be defined as:

```
#an LSTM layer with 5 memory cells/neurons
model = Sequential()
model.add(LSTM(5, input_shape=(10,1)))
...
```

The **number of samples does not have to be specified!**

Shuffling Samples (Not)



33

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

Samples within an epoch are usually shuffled. This is a good practice when working with MLP networks. However, if we are trying to **preserve state across samples**, then the **order of samples** in the training set **must be preserved**!

```
#an LSTM layer with 5 memory cells/neurons
model = Sequential()
model.add(LSTM(5, input_shape=(10,1)))
...
history = model.fit(X, y, shuffle=False, batch_size=10, epochs=100)
```

Once fit, an **history object** is returned, providing a **summary of the performance** of the model during training, recorded at each epoch. These metrics can be plotted and analyzed in regard to overfitting or underfitting (we will get back to this later!)

The Importance of Batches



34

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

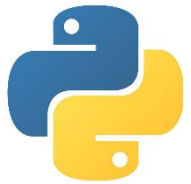
Mini-batch gradient descent with a **batch size of 32** is a common configuration for LSTMs.

We know that **each LSTM memory unit** maintains **internal state** that is accumulated! By **default**, the **internal state** of **all memory units** in the network is **reset after each batch!!** This is what is called a **Stateless LSTM** (default behaviour)!

This means that the **configuration of the batch size** imposes a relation between:

- The **efficiency** of **learning** (how many samples are processed before an update)
- The **speed** of **learning** (how often weights are updated)
- The **influence** of the **internal state** (how often internal state is reset)

LSTMs State Management



35

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

However, tf.Keras provides flexibility so that we may control **when to reset the internal state**! This is what is called a **Stateful LSTM**!

```
...  
#a batch size of 10 samples, with 5 timesteps and 1 feature  
model.add(LSTM(2, stateful=True, batch_input_shape=(10, 5, 1)))  
...
```

When stateful LSTM layers are used, we must also define the batch size as part of the input shape. The **batch_input_shape** argument requires a 3-dimensional tuple defined as **batch size, time steps** and **features**!

With **stateful LSTM** layers, the **internal state is shared across batches** (instead of being reset)!

LSTMs State Management



36

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

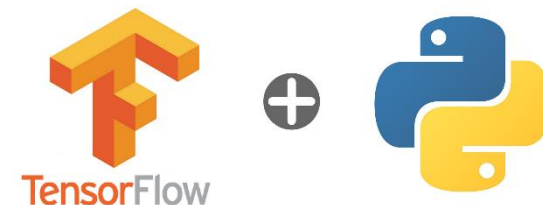
With **stateful LSTM** we have fine grained **control** over when to **reset the internal state**. Hence, if we want to reset it we must call the **reset_states()** function.

```
#reset the internal state at the end of each single epoch
for i in range(100):
    model.fit(X, y, epochs=1, batch_input_shape=(10, 5, 1))
    model.reset_states()
```

The **same batch size** used in the definition of the **stateful LSTM** must also be used **when making predictions**

```
...
predictions = model.predict(X, batch_size=10)
...
```


LSTMs State Management



37

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

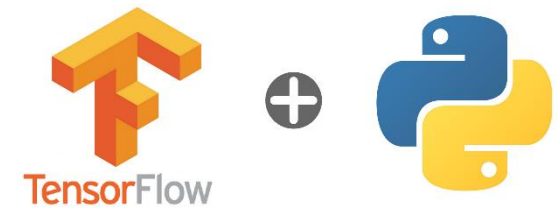
Hands On

The **internal state** is also accumulated when **evaluating** a network and when **making predictions**!

Therefore, if using a **stateful LSTM** we must reset the state after evaluating the network or after making predictions! **<- Important!!**

When sequences in different batches are related to each other, we should use stateful mode! Otherwise, when a sequence represents a complete sequence, we should go with stateless mode!

Some Last Tips



38

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

- LSTM **input must be 3D**. The meaning of the 3 input dimensions are:
 - Samples
 - Timesteps
 - Features
- The **LSTM input layer** is defined by the **input shape argument** on the first hidden layer
- The input shape argument takes **a tuple of two values** that define:
 - Timesteps
 - Features
- The **number of samples** is assumed to be **1 or more**
- The **reshape()** function on NumPy arrays can be used to reshape your **1D or 2D data to be 3D**

Vanilla LSTMs



39

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

The simplest LSTM configuration is the **Vanilla LSTM**

- We will name it Vanilla to differentiate it from deep LSTMs



The Echo Sequence Prediction Problem



40

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On



The Echo Sequence Prediction Problem



41

LSTMs

Seq. Prediction

Preparing Data

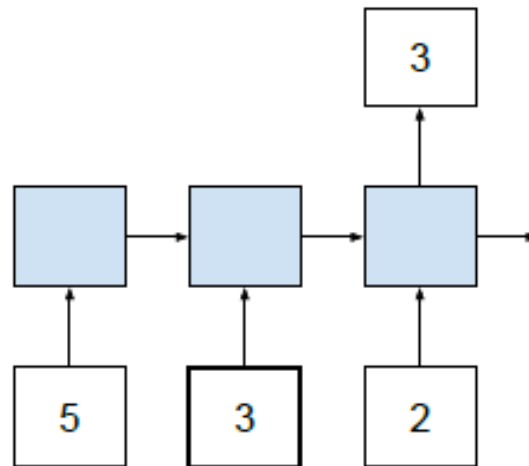
TF AND LSTMs

Hands On

A nice problem for demonstrating the **memory capability** of LSTMs!

The task is that, **given a sequence of random integers** as input, to **output the value of a random integer at a specific timestep** that is not specified to the model!

For example - given an input sequence of random integers [5, 3, 2] with the chosen timestep being the second value, then the expected output is 3 (technically, it is a sequence classification problem).



The Echo Sequence Prediction Problem



42

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

Lets first **generate random sequences** and then **one-hot encode such sequences!**

```
#generate a sequence of random numbers [0, n_features=10) with sequence length=5
def generate_sequence(length, nr_features):
    return [randint(0, nr_features-1) for _ in range(length)]

#one hot encoded sequence
def one_hot_encode(sequence, nr_features):
    encoded = list()
    for value in sequence:
        one_hot_encoded = np.zeros(nr_features)
        one_hot_encoded[value] = 1
        encoded.append(one_hot_encoded)
    return array(encoded)

#decode a one hot encoded sequence
def one_hot_decode(encoded_seq):
    return [argmax(value) for value in encoded_seq]
```

The Echo Sequence Prediction Problem



43

LSTMs

Seq. Prediction

Preparing Data

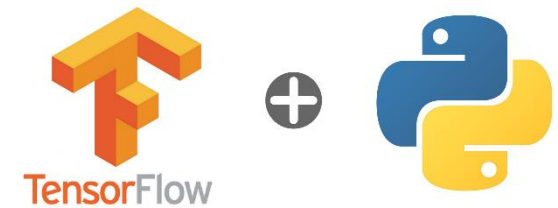
TF AND LSTMs

Hands On

Finally, **reshape** the **input sequence to be 3D** and **create a function** to create samples for us that are already encoded and that **returns the X and the y** (sequence as supervised)

```
#function to generate one sequence sample of random numbers [0, n_features=10)
#length of sequence (5) is the timesteps; the one hot encoded sequence (10) are the feature
#returns both the X (input) and the y (label)
def generate_sample(length, nr_features, out_index):
    #generate the sequence
    sequence = generate_sequence(length, nr_features)
    #one hot encode it
    encoded = one_hot_encode(sequence, nr_features)
    #reshape it to be 3D (1 sample, length timesteps, nr_features features)
    X = encoded.reshape((1, length, nr_features))
    #select the output (y) by getting the encoded value at the specified out_index
    #this must remain consistent for all generated examples for a model, so that it can learn
    y = encoded[out_index].reshape(1, nr_features)
    return X, y
```

The Echo Sequence Prediction Problem



44

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

Time to **define and compile a Stateless Model!** Lets use sequences with 5 timesteps, each one having a value between 0 and 9.

```
#control variables
timesteps = 5
features = 10
out_index = 2
epochs = 500

#define and compile the model
model = tf.keras.Sequential()
model.add(tf.keras.layers.LSTM(16, input_shape=(timesteps, features)))
model.add(tf.keras.layers.Dense(features, activation = 'softmax' ))
model.compile(
    loss= tf.keras.losses.categorical_crossentropy,
    optimizer= tf.keras.optimizers.Adam(learning_rate=0.001),
    metrics=['accuracy'])

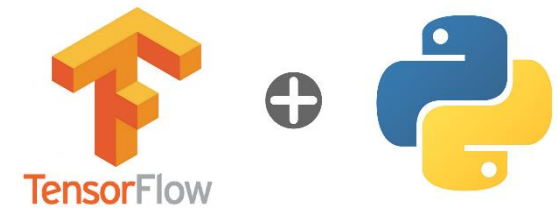
#printing the model's summary
print(model.summary())
```

Layer (type)	Output Shape	Param#
=====		
lstm (LSTM)	(None, 16)	1728

dense (Dense)	(None, 10)	170
=====		
Total params: 1,898		
Trainable params: 1,898		
Non-trainable params: 0		

None		

The Echo Sequence Prediction Problem



45

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

Now we **fit the Model**. Accuracy will be either 0 or 1 because we are making sequence classification prediction on one sample and reporting the result. We then **evaluate it**!

```
#fit the model
for i in range(epochs):
    X, y = generate_sample(timesteps, features, out_index)
    history = model.fit(X, y, shuffle=False, epochs=1, verbose=0)
    print('Epoch: %d; Loss: %.2f; Accuracy: %.2f' % (i,
    history.history['loss'][0], history.history['accuracy'][0]))

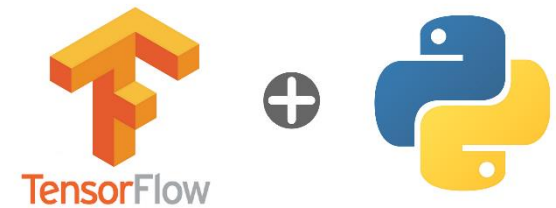
#evaluate the model by simply making predictions on randomly
#generated sequences and counting the number of correct predictions
correct = 0
for i in range(100):
    X, y = generate_sample(timesteps, features, out_index)
    yhat = model.predict(X)
    if one_hot_decode(yhat) == one_hot_decode(y):
        correct += 1
print('Accuracy: %.2f' % ((correct/100)*100.0))
```



```
...
Epoch: 4994; Loss: 0.20;
Accuracy: 1.00
Epoch: 4995; Loss: 1.84;
Accuracy: 0.00
Epoch: 4996; Loss: 0.37;
Accuracy: 1.00
Epoch: 4997; Loss: 1.15;
Accuracy: 1.00
Epoch: 4998; Loss: 0.49;
Accuracy: 1.00
...

Accuracy: 57.00
```

The Echo Sequence Prediction Problem



46

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

Now we **make predictions** on new randomly generated sequences (well, pretty much the same we did when evaluating)!

```
#prediction on new sequence data
```

```
X, y = generate_sample(timesteps, features, out_index)
```

```
yhat = model.predict(X)
```

```
#print values
```

```
print('Sequence: %s' % [one_hot_decode(x) for x in X])
```

```
print('Expected: %s' % one_hot_decode(y))
```

```
print('Predicted: %s' % one_hot_decode(yhat))
```

```
Sequence: [[5, 8, 7, 7, 5]]
```

```
Expected: [7]
```

```
Predicted: [7]
```



Deep LSTMs



47

LSTMs

Seq. Prediction

Preparing Data

TF AND LSTMs

Hands On

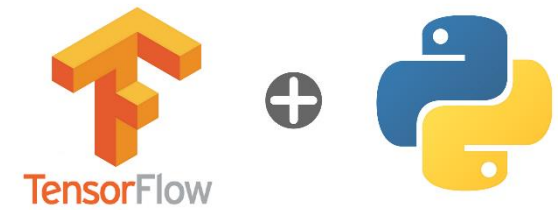
Long Short-Term Memory Networks for Traffic Flow Forecasting

15

Algorithm 3: Function for LSTM model's conception and compilation.

```
Input: timesteps, multisteps, features, h_layers=2, h_neurons=64,  
activation='relu', dropout_rate=0.5, deep_dense=False  
Output: Sequential LSTM Model  
model = Sequential()  
while  $i \in \text{range}(h\_layers)$  do  
    if  $i == 0$  then  
        if  $i+1 == h\_layers$  then  
            model.add(CuDNNLSTM(h_neurons, return_sequences=False,  
                                input_shape=(timesteps, features)))  
        else  
            model.add(CuDNNLSTM(int(h_neurons/2),  
                                return_sequences=True, input_shape=(timesteps, features)))  
            model.add(Dropout(dropout_rate))  
        end  
    else if  $i+1 == h\_layers$  then  
        model.add(CuDNNLSTM(h_neurons  $\times$  2, return_sequences=False))  
    else  
        model.add(CuDNNLSTM(h_neurons, return_sequences=True))  
        model.add(Dropout(dropout_rate))  
    end  
end  
model.add(Dense(h_neurons, activation=activation))  
model.add(Dropout(dropout_rate))  
if deep_dense then  
    model.add(Dense(int(h_neurons/2), activation=activation))  
    model.add(Dropout(dropout_rate))  
model.add(Dense(multisteps))  
model.compile(loss=rmse, optimizer= Adam(), metrics = [mae, rmse])  
return model
```

Glossary



48

LSTMs

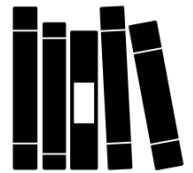
Seq. Prediction

Preparing Data

TF and LSTMs

HANDS ON

- **Gates**: used by RNNs such as LSTMs or GRUs, to decide how to update/handle the internal state of a memory cell
- **LSTMs**: Long Short-Term Memory networks, a type of RNNs
- **Input Shape for LSTMs**: 3D vector of (samples, timesteps, features)
- **Single-step vs Multi-step**: forecasting just the next value of the sequence (ex.: the next hour) vs forecasting multiple timesteps beyond the input sequences (ex.: the next twelve hours)
- **Stateless vs Stateful LSTMs**: internal state of all memory units in the network is reset after each batch automatically (default behavior) vs controlling when to reset the internal state
- **Univariate vs Multivariate**: using a single feature as input (ex.: speed_diff) vs using multiple features as input (ex.: speed_diff, time_diff and temperature)



Resources

49

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

HANDS ON

- Official Documentation
 - https://www.tensorflow.org/api_docs/python/tf/keras/layers/LSTM
 - <https://www.tensorflow.org/guide/keras/rnn>
 - ...
- Papers, Books, online courses, tutorials...
 - Understanding LSTM Networks
<https://colah.github.io/posts/2015-08-Understanding-LSTMs>
 - The Unreasonable Effectiveness of Recurrent Neural Networks
<http://karpathy.github.io/2015/05/21/rnn-effectiveness>
 - Learning to Forget: Continual Prediction with LSTM, 1999
<https://pdfs.semanticscholar.org/e10f/98b86797ebf6c8caea6f54cacbc5a50e8b34.pdf>
 - Hochreiter, S. & Schmidhuber, J., “Long Short-Term Memory”, Neural Computation 9(8), pp. 1735-1780, 1997. <http://www.bioinf.jku.at/publications/older/2604.pdf>
 - (Book) Long Short-Term Memory Networks With Python by Jason Brownlee

Hands On



50

LSTMs

Seq. Prediction

Preparing Data

TF and LSTMs

HANDS ON

Spyder (Python 3.6)

File Edit Search Source Run Debug Consoles Projects Tools View Help

Editor - C:\data\PythonWorkspace\dev\meanshift_algorithm.py

```
37 class Mean_Shift:
38     def __init__(self, radius=None, radius_normalize_step = 150):
39         self.radius = radius
40         self.radius_normalize_step = radius_normalize_step
41
42     def fit(self, data):
43
44         if self.radius == None:
45             all_data_centroid = np.average(data, axis=0)
46             all_data_norm = np.linalg.norm(all_data_centroid)
47             self.radius = all_data_norm/self.radius_normalize_step
48
49         centroids = {}
50
51         #initialize centroids
52         for i in range(len(data)):
53             centroids[i] = data[i]
54
55         weights = [1 for i in range(self.radius_normalize_step)]
56
57         while True:
58             new_centroids = []
59             for i in centroids:
60                 in_range = []
61                 centroid = centroids[i]
62
63                 for featureset in data:
64                     distance = np.linalg.norm(featureset-centroid)
65                     if distance == 0:
66                         distance = 0.0000000001
67                     weight_index = int(distance/self.radius)
68                     if weight_index > self.radius_normalize_step-1:
69                         weight_index = self.radius_normalize_step-1
70                     to_add = (weights[weight_index]**2)*[featureset]
71                     in_range += to_add
72
73             new_centroid = np.average(in_range, axis=0)
```

Variable explorer

Name	Type	Size	Value
batch_size	int	1	100
mnist	contrib.learn.python.learn.datasets.base.Datasets	3	Datasets object of...
n_classes	int	1	10
n_nodes_hl1	int	1	500
n_nodes_hl2	int	1	500
n_nodes_hl3	int	1	500

Python console

See 'tf.nn.softmax_cross_entropy_with_logits_v2'.

Epoch 0 completed out of 10 loss: 1666037.4677734375
Epoch 1 completed out of 10 loss: 377809.3128890991
Epoch 2 completed out of 10 loss: 201302.4857263565
Epoch 3 completed out of 10 loss: 119427.91378033161
Epoch 4 completed out of 10 loss: 72651.25679710507
Epoch 5 completed out of 10 loss: 45327.621502393486
Epoch 6 completed out of 10 loss: 31955.17812934518
Epoch 7 completed out of 10 loss: 23664.35610633137
Epoch 8 completed out of 10 loss: 18248.740643078025
Epoch 9 completed out of 10 loss: 19962.00065876091
Accuracy: 0.9511

In [2]:

Python console History log